

Crop Recommendation for India

Tasotti Kerstin

2022-03-08

1. INTRODUCTION

I decided to analyze a crop recommendation dataset as my own submission project for the edX course HarvardX PH125.9x - Data Science:Capstone Course. The aim of this project is to demonstrate the acquired skills in R programming and their analysis in a real world datasets. This dataset is to recommend optimum crops to be cultivated by farmers based on several parameters and help them make an informed decision before cultivation. The data used in this project is loaded from www.kaggle.com. and contains data from the Indian agriculture. This data is relatively simple with only 7 soil and environmental conditions. The data contains Nitrogen, Phosphorus, Potassium and pH values of the soil. The environmental conditions are the humidity, temperature and rainfall which are collected for a particular crop. In the dataset is no information on the harvest yield. Precision agriculture is in trend nowadays. It helps the farmers to get informed decision about the farming strategy. Here, I present you a dataset which would allow the users to build a predictive model to recommend the most suitable crops to grow in a particular farm based on various parameters.

Data fields: Soil condition N - ratio of Nitrogen content in soil in ppm P - ratio of Phosphorus content in soil in ppm K - ratio of Potassium content in soil in ppm
ph - pH value of the soil

Environment condition temperature - temperature in degrees Celsius humidity - relative humidity in % rainfall - rainfall in mm

A recommendation system is a subclass of information filtering system that seeks to predict the “rating” or “preference” a user would give to an item. Recommendation systems is one of the most used machine learning algorithms and will be used in nearly all different areas of our life (Trading, Hospitality, Travelling,...). Companies like Amazon use these systems to learn more about their customer and provide them with products more effectively.

This dataset is prepared, and different analysis were done to develop an algorithm of a machine learning which can predict crop rate and calculate the RMSE (Root mean square error). The RMSE is a KPI to measure the differences between the predicted values of a model and the actual values seen in the data. Therefore the dataset is split into a training set (train) and a final hold-out test set (validation). The objective was for the final algorithm to predict ratings with a root mean square error (RMSE) less as possible. We want to generate a nearly accurate machine learning algorithm.

2. DATASET

For this project we focus on the Kaggle dataset collected by Atharva Ingle and it can be found on Kaggle website (<https://www.kaggle.com/atharvaingle/crop-recommendation-dataset> (<https://www.kaggle.com/atharvaingle/crop-recommendation-dataset>)). The data is uploaded in GitHub and loaded into the model.

2.1. DATALOAD

```
`#####`

# Load of full, train set, validation set (final hold-out test set)

`#####`

if(!require(tidyverse)) install.packages("tidyverse", repos = "http://cran.us.r-project.org")
if(!require(ggplot2)) install.packages("ggplot2", repos = "http://cran.us.r-project.org")
if(!require(Metrics)) install.packages("Metrics", repos = "http://cran.us.r-project.org")
if(!require(caret)) install.packages("caret", repos = "http://cran.us.r-project.org")
if(!require(data.table)) install.packages("data.table", repos = "http://cran.us.r-project.org")
if(!require(readr)) install.packages("readr", repos = "http://cran.us.r-project.org")
if(!require(tinytex)) install.packages("tinytex", repos = "http://cran.us.r-project.org")
if(!require(knitr)) install.packages("knitr", repos = "http://cran.us.r-project.org")

library(tidyverse)
library(ggplot2)
library(Metrics)
library(caret)
library(data.table)
library(readr)
library(tinytex)
library(knitr)

## Crop dataset:

### <https://www.kaggle.com/atharvaingle/crop-recommendation-dataset>

### <https://www.kaggle.com/atharvaingle/crop-recommendation-dataset/download>

# Crop dataset:
# https://github.com/KerstinTasotti/Crop.git

###Full dataset:
Crop_rec <- read.csv("https://raw.githubusercontent.com/KerstinTasotti/Crop/main/Crop_recommendation.csv", header=TRUE, stringsAsFactors = FALSE)

### Validation set will be 10% of Crop data
set.seed(1, sample.kind="Rounding")
train_index <- createDataPartition(y=Crop_rec$label, p=0.9, list=FALSE, time=1)
train<- Crop_rec[train_index,]
validation <- Crop_rec[-train_index,]
```

3. ANALYSIS OF THE DATA

All analysis in this section will be done with the training set (train). The validation set will be used for the final test of the developed algorithm. The dataset is shuffled and the first 1980 lines of the 2200 lines will be taken for the training set.

At first the structure of the dataset will be analyzed to get familiar with it. The data set contains 2200 rows and 8 variables (N, P, K, temperature, humidity, ph, rainfall, label). The train set contains 1980 rows and the same 8 variables.

```
str(train)

head(train) %>% print.data.frame
```

Each label has 90 entries. Therefore, we select for further calculations $m=90$.

```
train %>% group_by(label) %>% summarise(count=n()) %>% arrange(desc(count))
m <- 90
```

In the summary we can see the min, max and mean values of the different ground condition.

```
summary(train)
```

N		P		K		temperature	
Min.	: 0.00	Min.	: 5.00	Min.	: 5.00	Min.	: 8.826
1st Qu.:	21.00	1st Qu.:	28.00	1st Qu.:	20.00	1st Qu.:	22.782
Median	: 37.00	Median	: 51.00	Median	: 32.00	Median	: 25.629
Mean	: 50.52	Mean	: 53.36	Mean	: 48.14	Mean	: 25.670
3rd Qu.:	84.00	3rd Qu.:	68.00	3rd Qu.:	49.00	3rd Qu.:	28.614
Max.	: 140.00	Max.	: 145.00	Max.	: 205.00	Max.	: 43.675

humidity		ph		rainfall		label	
Min.	: 14.26	Min.	: 3.505	Min.	: 20.36	Length: 1980	
1st Qu.:	60.37	1st Qu.:	5.969	1st Qu.:	64.51	Class : character	
Median	: 80.47	Median	: 6.420	Median	: 95.03	Mode : character	
Mean	: 71.50	Mean	: 6.468	Mean	: 103.49		
3rd Qu.:	89.92	3rd Qu.:	6.921	3rd Qu.:	124.22		
Max.	: 99.98	Max.	: 9.935	Max.	: 298.56		

There are 22 different labels in the dataset:

```
unique(train$label)
```

```
[1] "rice"      "maize"      "chickpea"   "kidneybeans" "pigeonpeas" "mothbeans"  "mungbean"
[2] "blackgram" "lentil"     "pomegranate"
[11] "banana"    "mango"      "grapes"     "watermelon"  "muskmelon"   "apple"      "orange"
[21] "papaya"    "coconut"    "cotton"
[21] "jute"      "coffee"
```

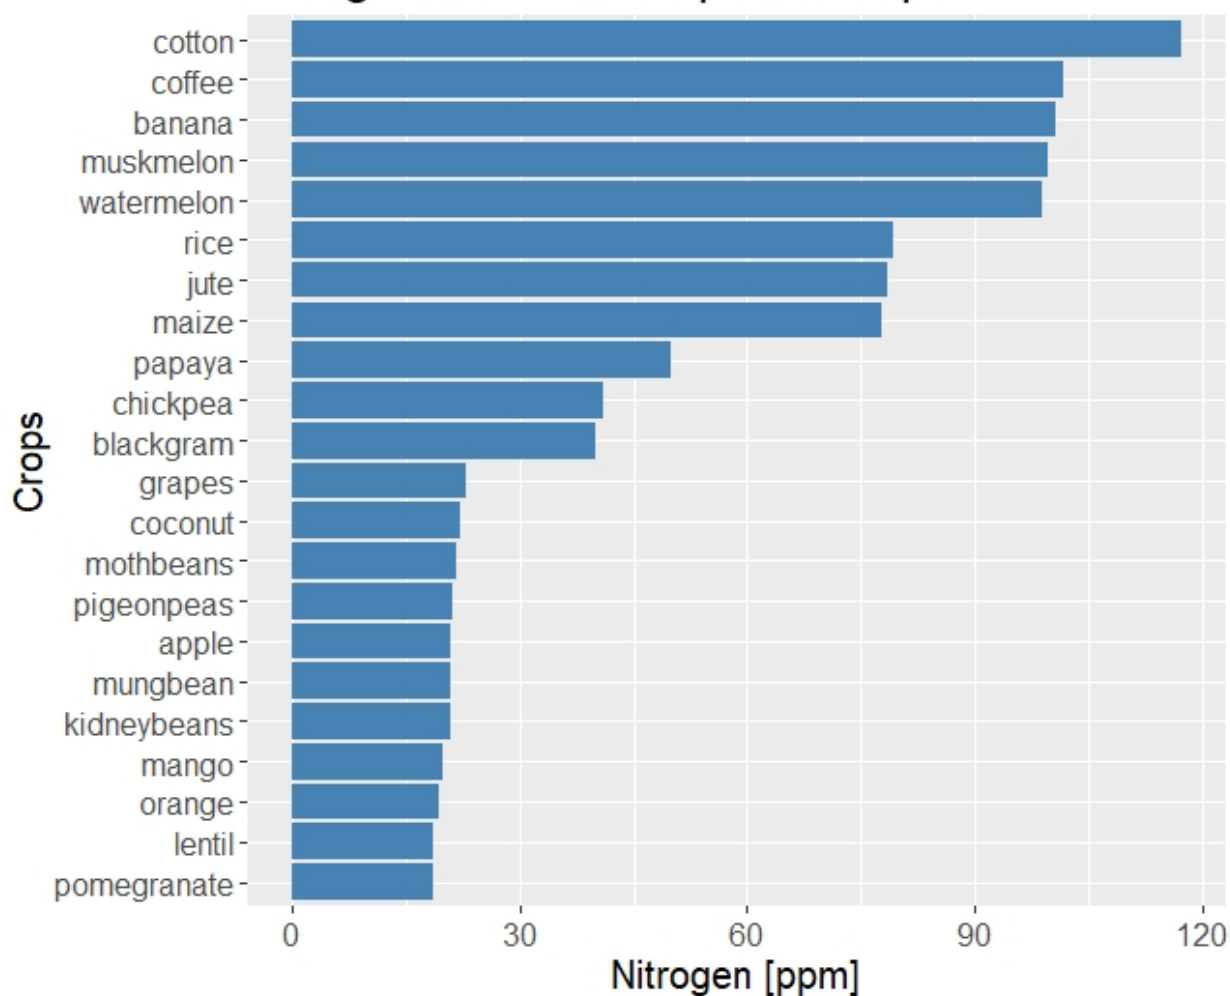
Nitrogen

Nitrogen is considered the most important component for supporting plant growth. Nitrogen is part of the chlorophyll molecule, which gives plants their green color and is involved in creating food for the plant through photosynthesis. Lack of nitrogen shows up as general yellowing (chlorosis) of the plant. The mean nitrogen value is 50.5ppm, whereby cotton and coffee need the most nitrogen and lentil and pomegranate need the lowest nitrogen value. A list of all mean Nitrogen values is summarized in the following table and bar chart:

```
train %>%group_by(label)%>%summarise(mean(N)) %>% print.data.frame()
```

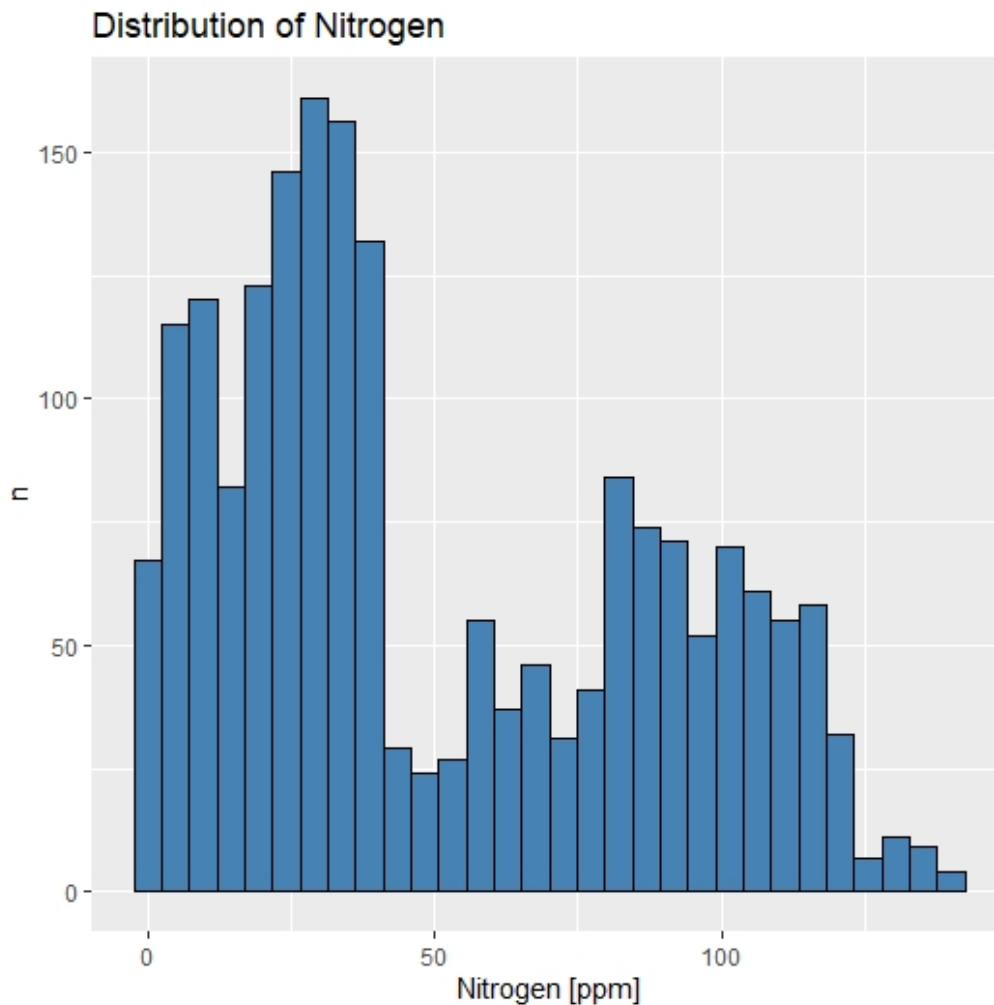
	label	mean(N)
1	apple	21.00000
2	banana	100.78889
3	blackgram	39.95556
4	chickpea	40.92222
5	coconut	22.06667
6	coffee	101.64444
7	cotton	117.28889
8	grapes	22.96667
9	jute	78.40000
10	kidneybeans	20.88889
11	lentil	18.67778
12	maize	77.75556
13	mango	19.84444
14	mothbeans	21.72222
15	mungbean	20.88889
16	muskmelon	99.74444
17	orange	19.42222
18	papaya	49.82222
19	pigeonpeas	21.11111
20	pomegranate	18.64444
21	rice	79.17778
22	watermelon	98.78889

Nitrogen content for optimal crops

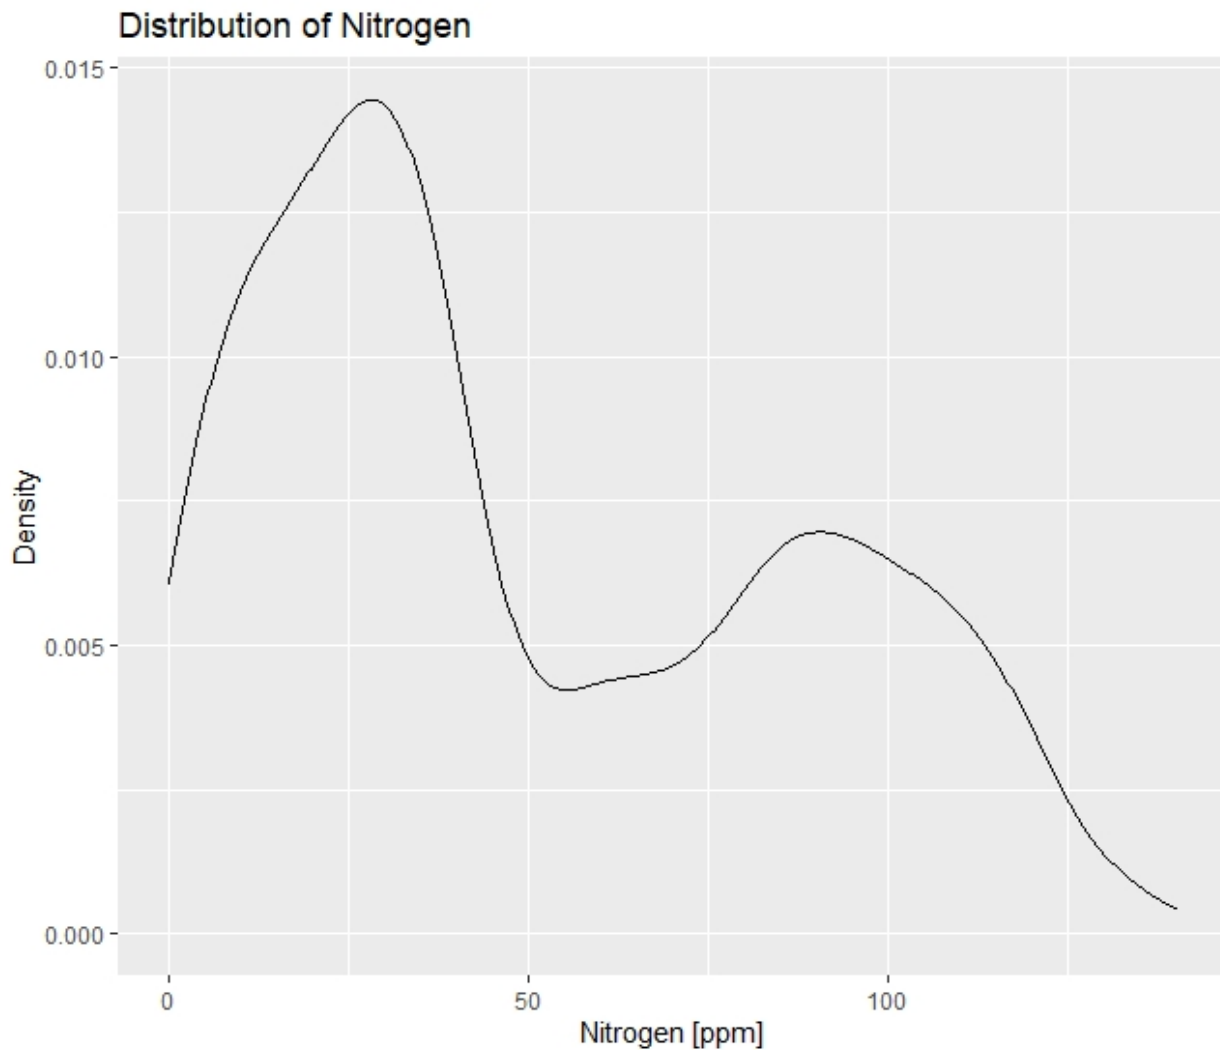


```
train %>% group_by(label) %>% ggplot(aes(x=reorder(label,N/m),y=N/m)) + geom_bar(stat = "identity", fill="steelblue") +labs(x="Crops", y="Nitrogen [ppm]", title="Nitrogen content for optimal crops") + theme(text = element_text(size=15))+coord_flip()
```

The Histogram and Density chart shows the distribution of Nitrogen. The most crops prefer a soil with lower Nitrogen content.



```
ggplot(train,aes(N)) +geom_histogram(bins=30, fill="steelblue",color="black")+labs(x="Nitrogen [ppm]",y="n")+ggtitle("Distribution of Nitrogen")
```



```
ggplot(train,aes(N)) +geom_density()+labs(x="Nitrogen [ppm]",y="Density")+ggtitle("Distribution of Nitrogen")
```

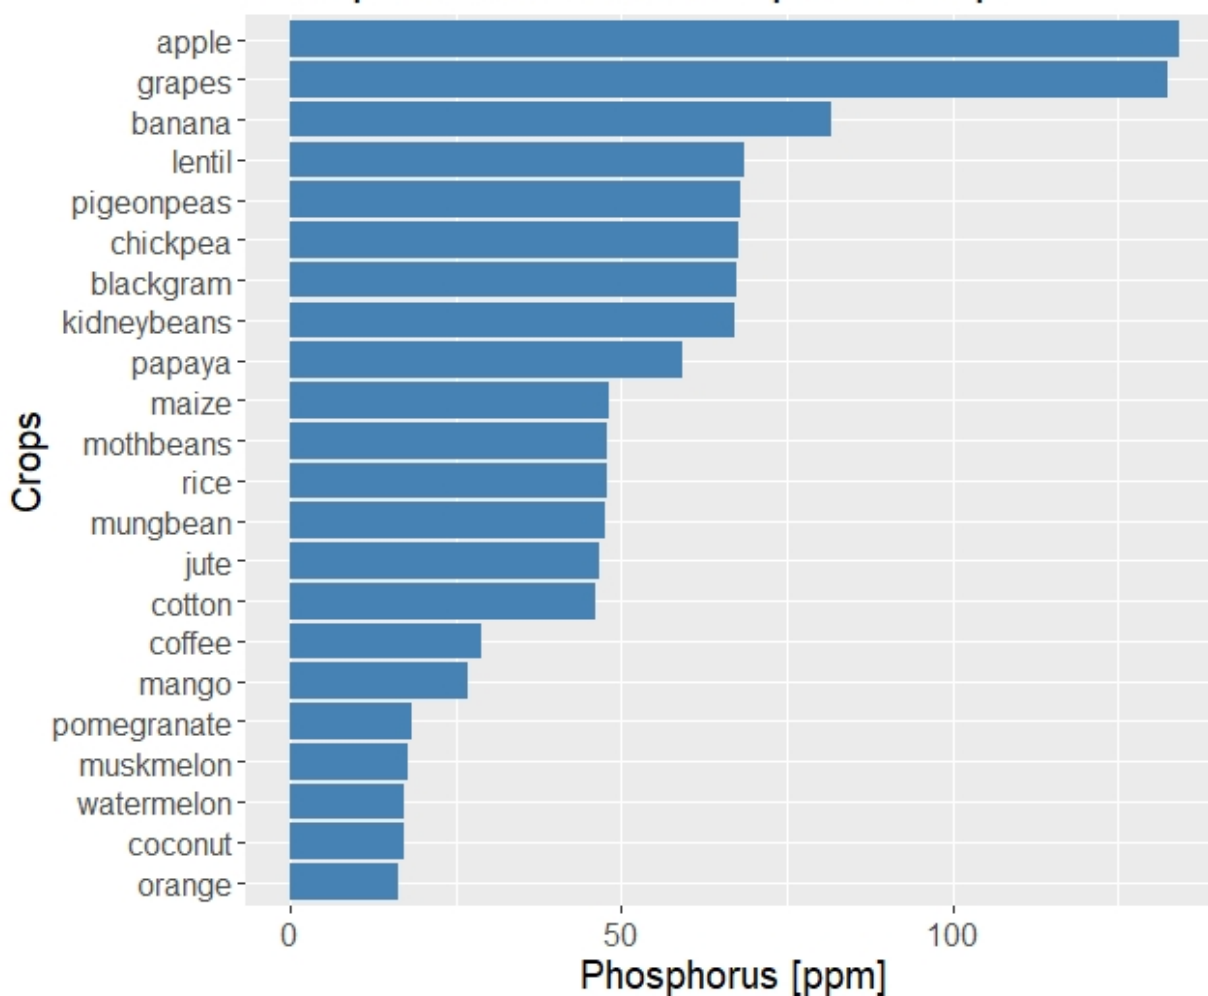
Phosphorus

Phosphorus is, therefore, important in cell division and development of new tissue. Phosphorus is also associated with complex energy transformations in the plant. Adding phosphorus to soil low in available phosphorus promotes root growth and winter hardiness, stimulates tillering, and often hastens maturity. Apple and grapes need ground with the highest phosphorus content and orange and coconut needs the lowest content. A list of all mean Phosphorus values is summarized in the following table and bar chart:

```
train %>%group_by(label)%>%summarise(mean(P)) %>% print.data.frame()
```

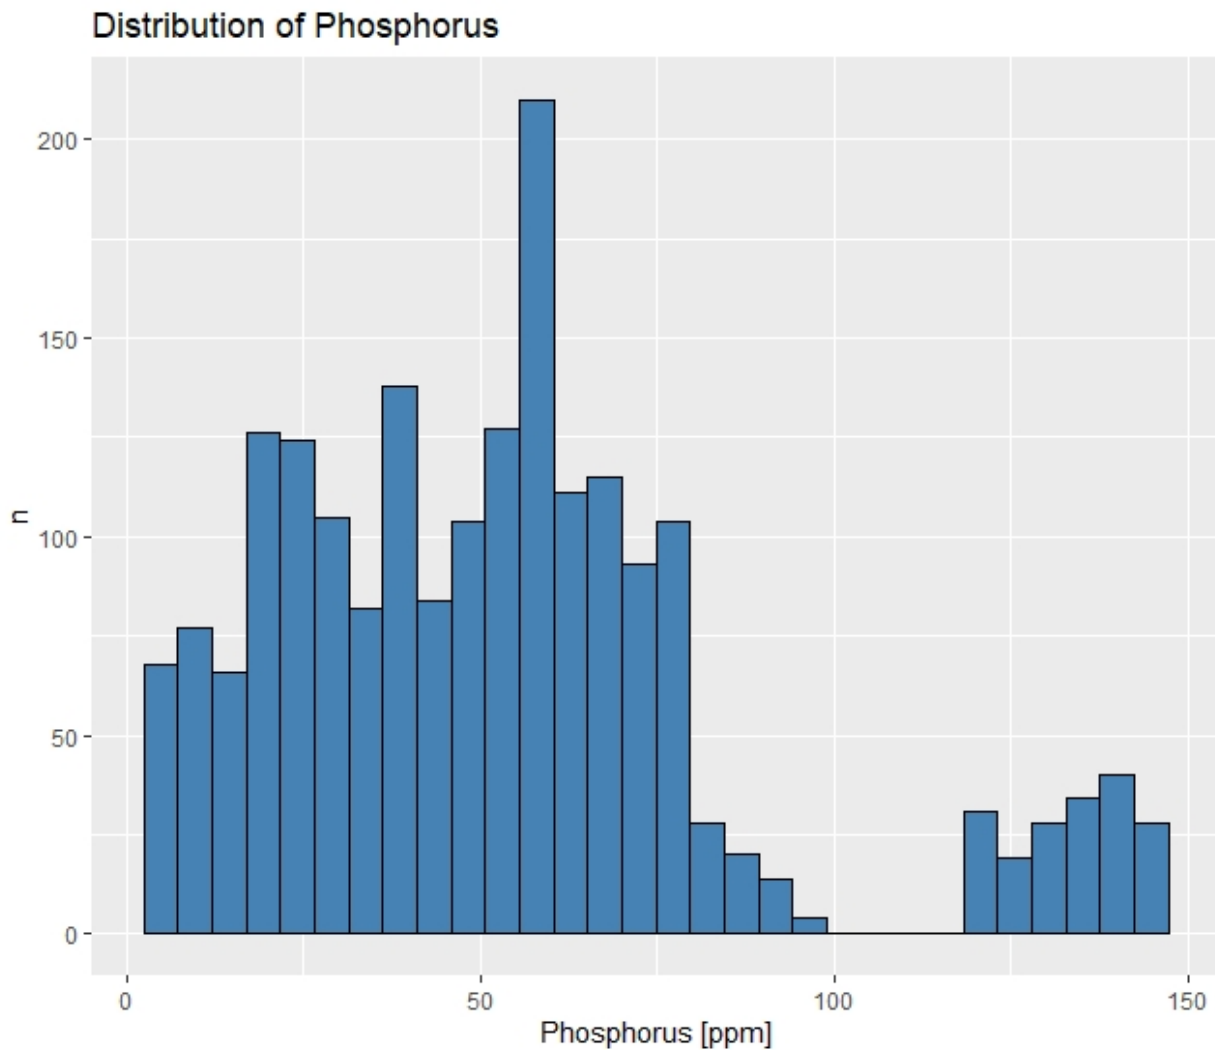
	label	mean(P)
1	apple	134.34444
2	banana	81.87778
3	blackgram	67.46667
4	chickpea	67.66667
5	coconut	17.11111
6	coffee	28.78889
7	cotton	46.18889
8	grapes	132.45556
9	jute	46.82222
10	kidneybeans	67.15556
11	lentil	68.54444
12	maize	48.17778
13	mango	26.86667
14	mothbeans	47.92222
15	mungbean	47.64444
16	muskmelon	17.72222
17	orange	16.45556
18	papaya	59.31111
19	pigeonpeas	67.93333
20	pomegranate	18.47778
21	rice	47.85556
22	watermelon	17.22222

Phosphorus content for optimal crops

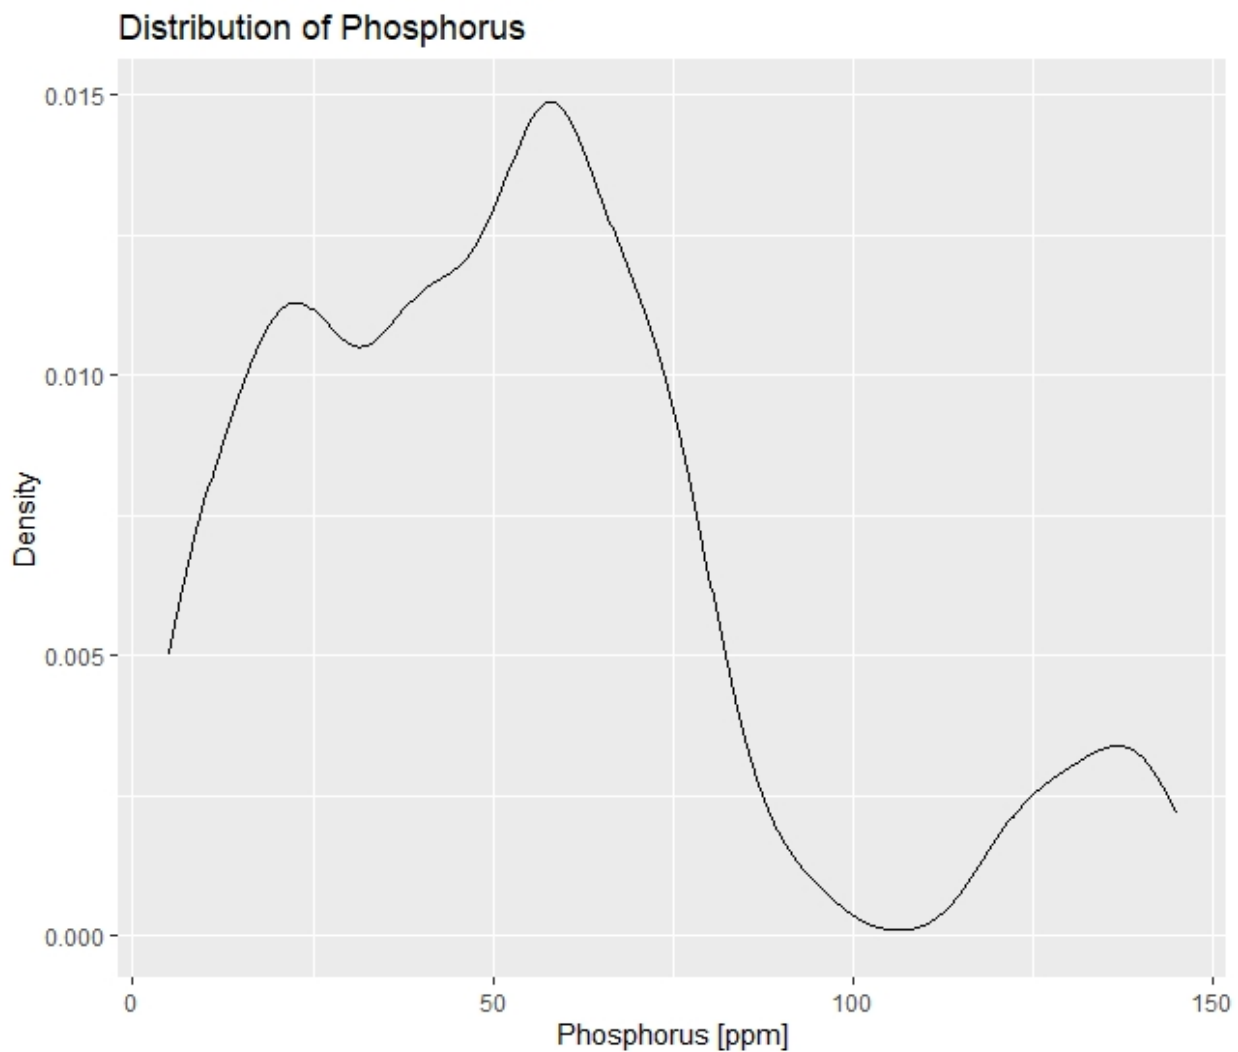


```
train %>% group_by(label) %>% ggplot(aes(x=reorder(label,P/m),y=P/m)) + geom_bar(stat = "identity", fill="steelblue") + labs(x="Crops", y="Phosphorus [ppm]", title="Phosphorus content for optimal crops") + theme(text = element_text(size=15))+coord_flip()
```

The Histogram and Density chart shows the distribution of Phosphorus. Only grapes and apples prefer a soil with a Phosphorus content higher than 100ppm.



```
ggplot(train,aes(P)) +geom_histogram(bins=30, fill="steelblue",color="black")+labs(x="Phosphorus [ppm]",y="n")+ggtitle("Distribution of Phosphorus")
```

```
ggplot(train,aes(P)) +geom_density()+labs(x="Phosphorus [ppm]",y="Density")+ggtitle("Distribution of Phosphorus")
```

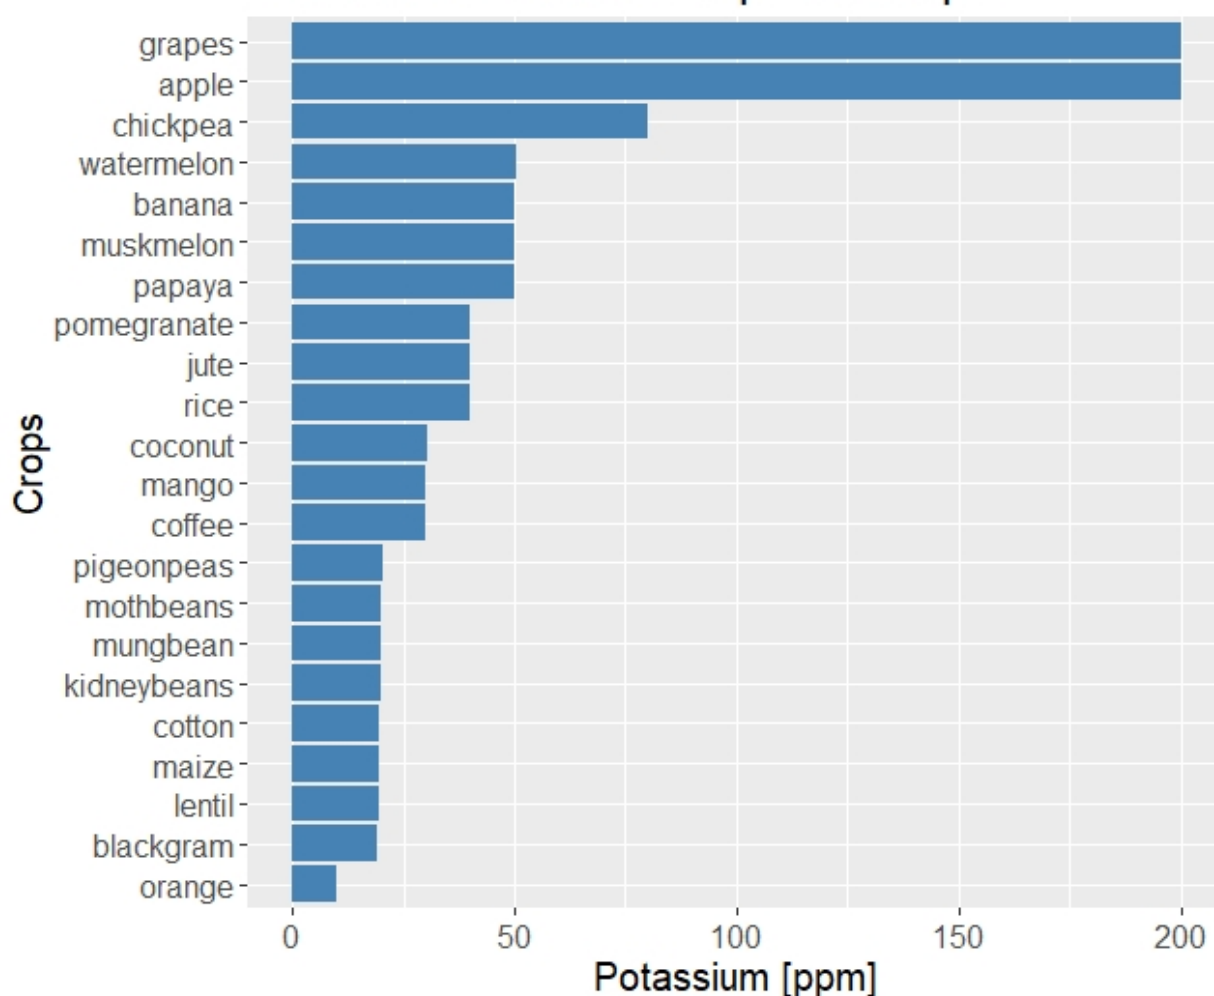
Potassium

Potassium helps photosynthesis, the process through which the sugars and energy that the plant needs for its development are formed and converted. Potassium also controls the opening and closing of the leaf stomata, which regulate the water status in the plant. Grapes and apples need soil with very high content of potassium, whereby again orange needs a very low potassium content followed by lentil. A list of all mean Potassium values is summarized in the following table and bar chart:

```
train %>%group_by(label)%>%summarise(mean(K)) %>% print.data.frame()
```

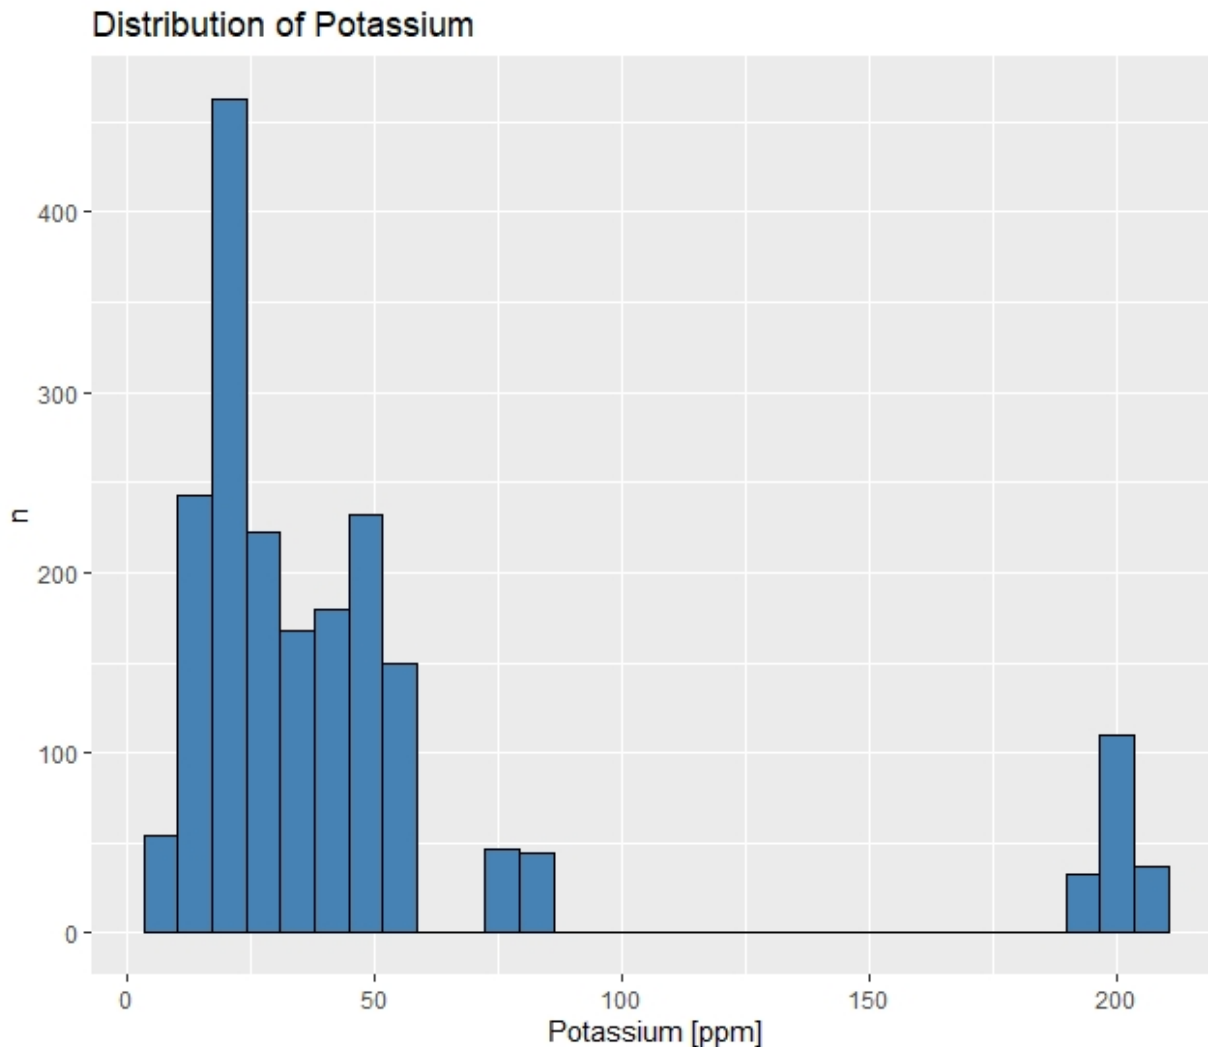
	label	mean(K)
1	apple	200.04444
2	banana	50.04444
3	blackgram	19.20000
4	chickpea	80.02222
5	coconut	30.55556
6	coffee	29.92222
7	cotton	19.57778
8	grapes	200.18889
9	jute	39.93333
10	kidneybeans	19.95556
11	lentil	19.41111
12	maize	19.53333
13	mango	29.92222
14	mothbeans	20.12222
15	mungbean	19.98889
16	muskmelon	50.03333
17	orange	10.01111
18	papaya	49.91111
19	pigeonpeas	20.41111
20	pomegranate	40.12222
21	rice	39.84444
22	watermelon	50.37778

Potassium content for optimal crops

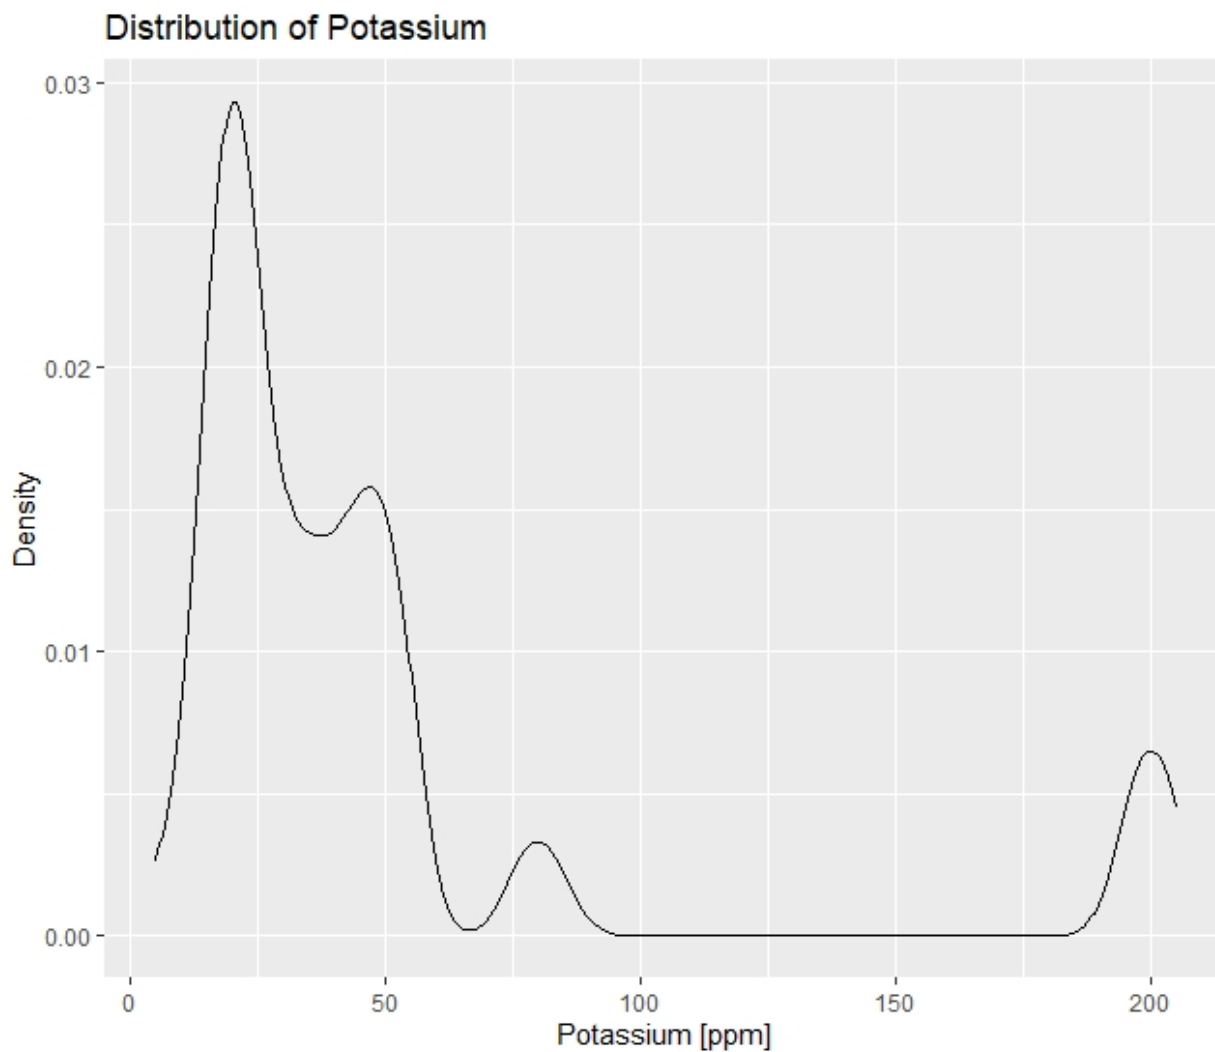


```
train %>% group_by(label) %>% ggplot(aes(x=reorder(label,K/m),y=K/m)) + geom_bar(stat = "identity", fill="steelblue") + labs(x="Crops", y="Potassium [ppm]", title="Potassium content for optimal crops") + theme(text = element_text(size=15))+coord_flip()
```

The Histogram and Density chart shows the distribution of Potassium. Just as with phosphorus content, grapes and apple require very high levels of potassium in the soil (around 200ppm). In contrast, all other crops require a potassium content of less than 85ppm.



```
ggplot(train,aes(K)) +geom_histogram(bins=30, fill="steelblue",color="black")+labs(x="Potassium [ppm]",y="n")+ggtitle("Distribution of Potassium")
```



```
ggplot(train,aes(K)) +geom_density()+labs(x="Potassium [ppm]",y="Density")+ggtitle("Distribution of Potassium")
```

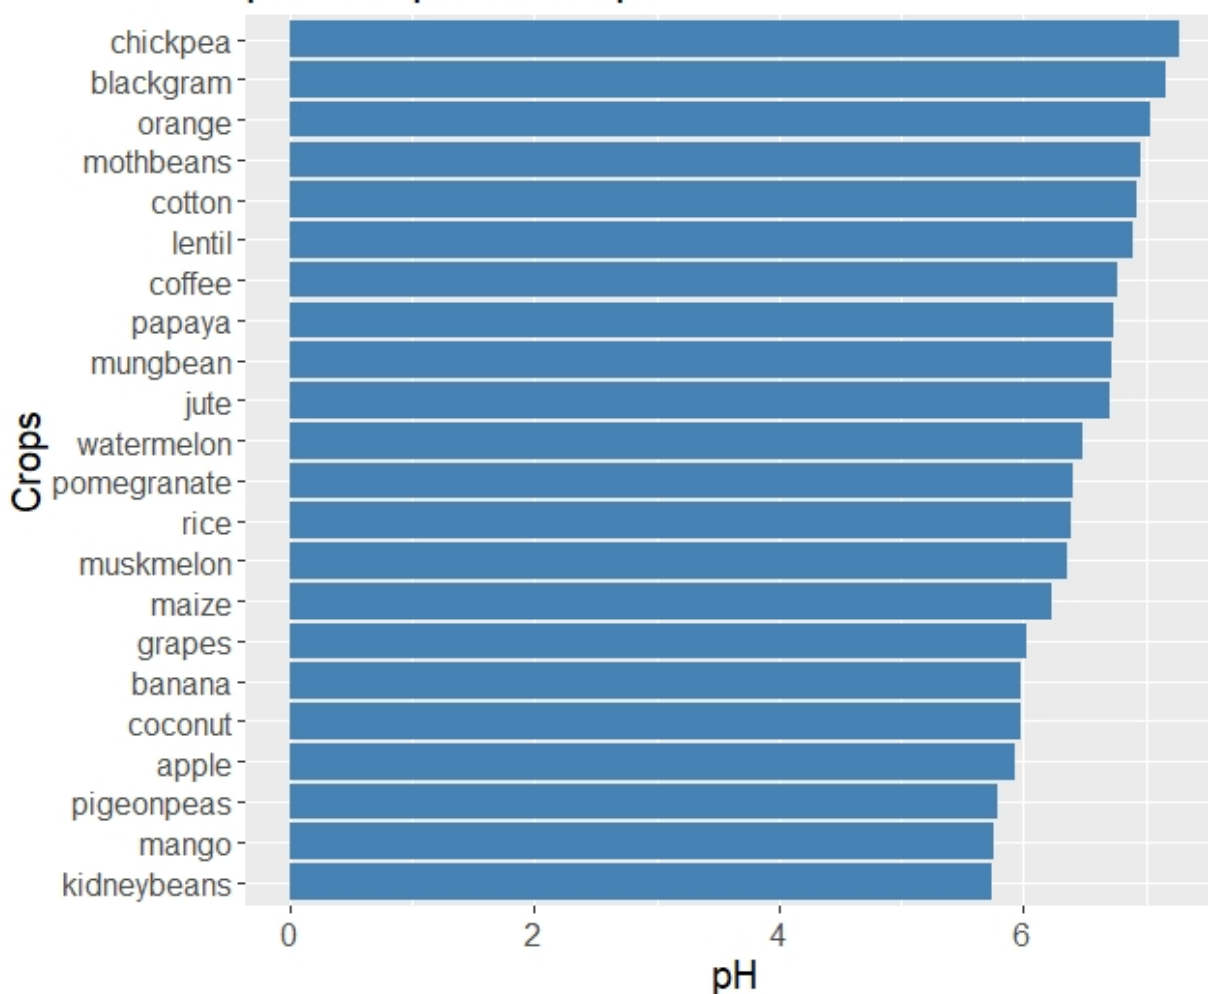
pH

In general, pH values between 6 and 7.5 are optimum for crop and forage production and nutrient uptake. Soil pH impacts nutrient availability and overall soil health. Soil acidification can be an indication of excessive application of nitrogen fertilizer. The range of pH is bigger than the optimal pH range should be in general, and the mean values of the different labels are between 5.74 and 7.27. Chickpea and blackgram need a neutral soil and papaya, mango and kidneybeans need a light acid soil with a pH of about 5.75. A list of all mean pH values is summarized in the following table and bar chart:

```
train %>%group_by(label)%>%summarise(mean(ph)) %>% print.data.frame()
```

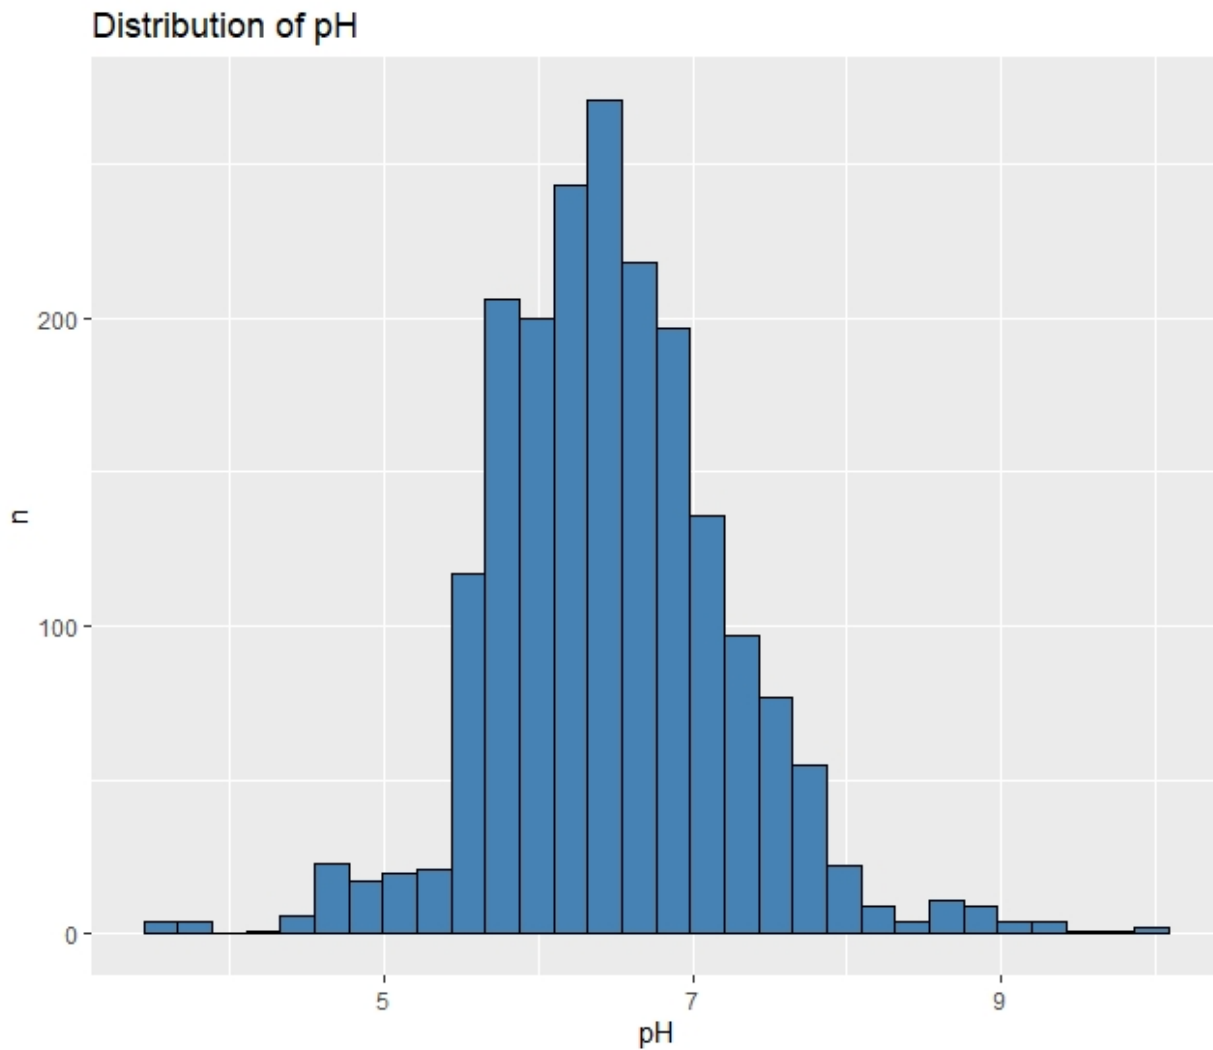
	label	mean(ph)
1	apple	5.936633
2	banana	5.980508
3	blackgram	7.159371
4	chickpea	7.278753
5	coconut	5.973727
6	coffee	6.771006
7	cotton	6.924236
8	grapes	6.021094
9	jute	6.714211
10	kidneybeans	5.748075
11	lentil	6.902046
12	maize	6.231523
13	mango	5.764537
14	mothbeans	6.963124
15	mungbean	6.727810
16	muskmelon	6.363348
17	orange	7.038504
18	papaya	6.735134
19	pigeonpeas	5.781993
20	pomegranate	6.413503
21	rice	6.391305
22	watermelon	6.486128

pH for optimal crops

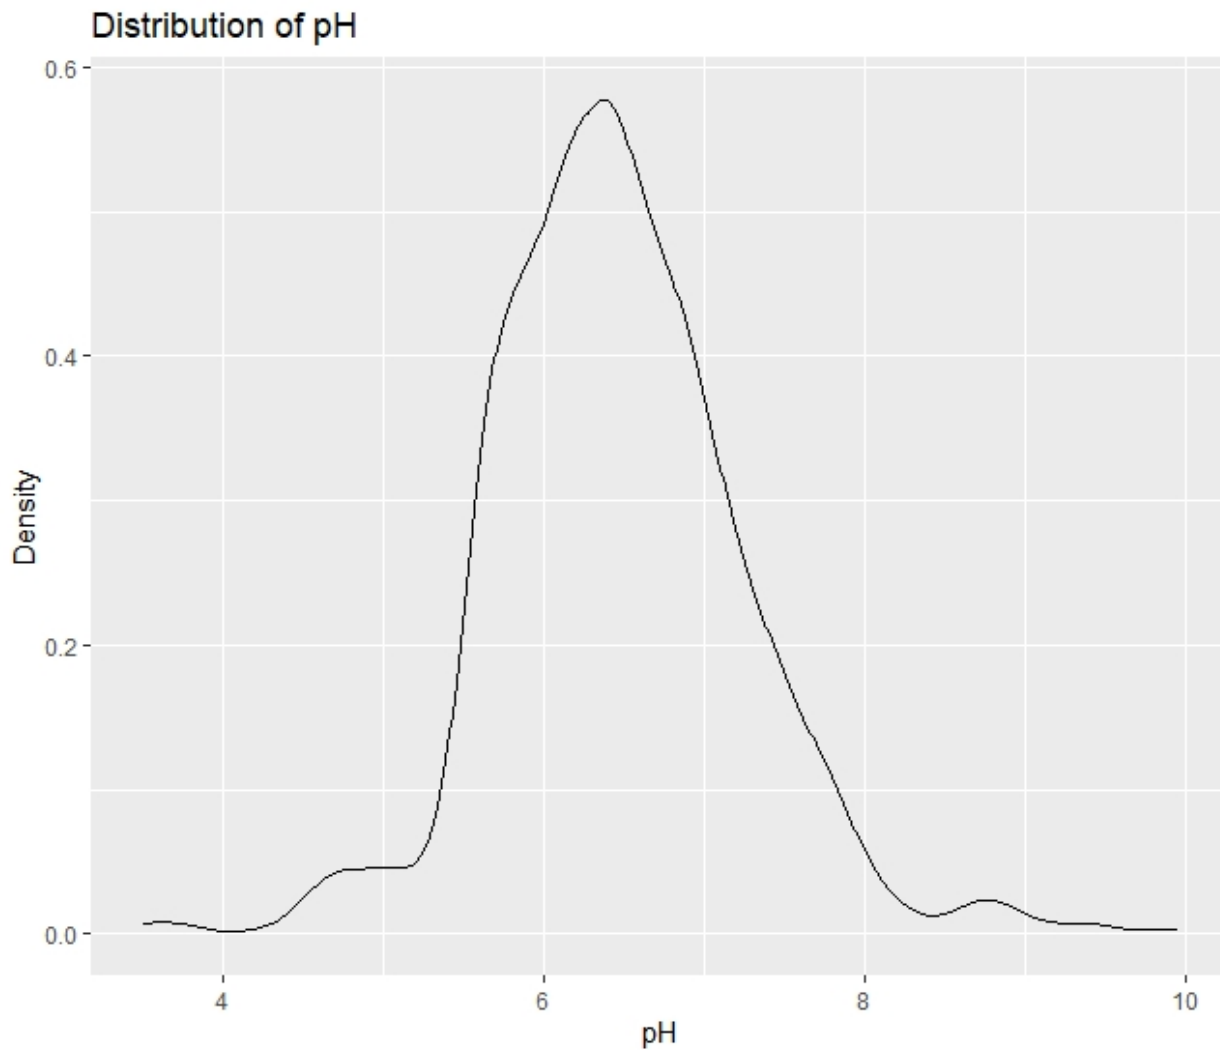


```
train %>% group_by(label) %>% ggplot(aes(x=reorder(label,ph/m),y=ph/m)) + geom_bar(stat = "identity", fill="steelblue") +labs(x="Crops", y="pH", title="pH for optimal crops") + theme(text = element_text(size=15))+coord_flip()
```

The Histogram and Density chart shows the distribution of pH. The pH range is normal distributed with a mean value of 6.47.



```
ggplot(train,aes(ph)) +geom_histogram(bins=30, fill="steelblue",color="black")+labs(x="pH",y="n")+ggtitle("Distribution of pH")
```



```
ggplot(train,aes(ph)) +geom_density()+labs(x="pH",y="Density")+ggtitle("Distribution of pH")
```

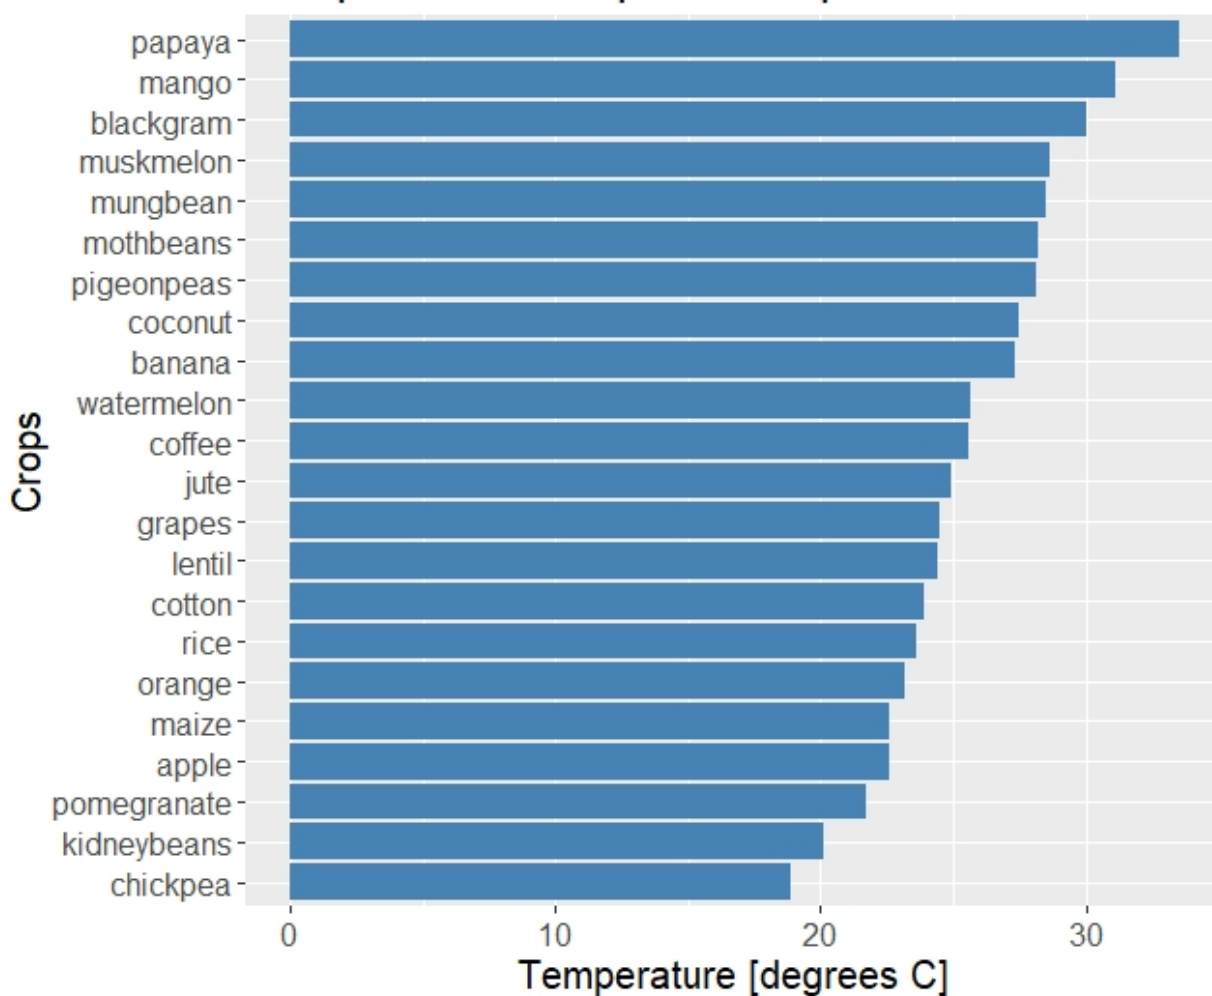
Temperature

As temperature increases (up to a point), photosynthesis, transpiration, and respiration increase. When combined with day-length, temperature also affects the change from vegetative (leafy) to reproductive (flowering) growth. The range of different temperature is between 18.86 (chickpea) and 30.02 (banana) degrees Celsius. A list of all mean Temperature values is summarized in the following table and bar chart:

```
train %>%group_by(label)%>%summarise(mean(temperature)) %>% print.data.frame()
```

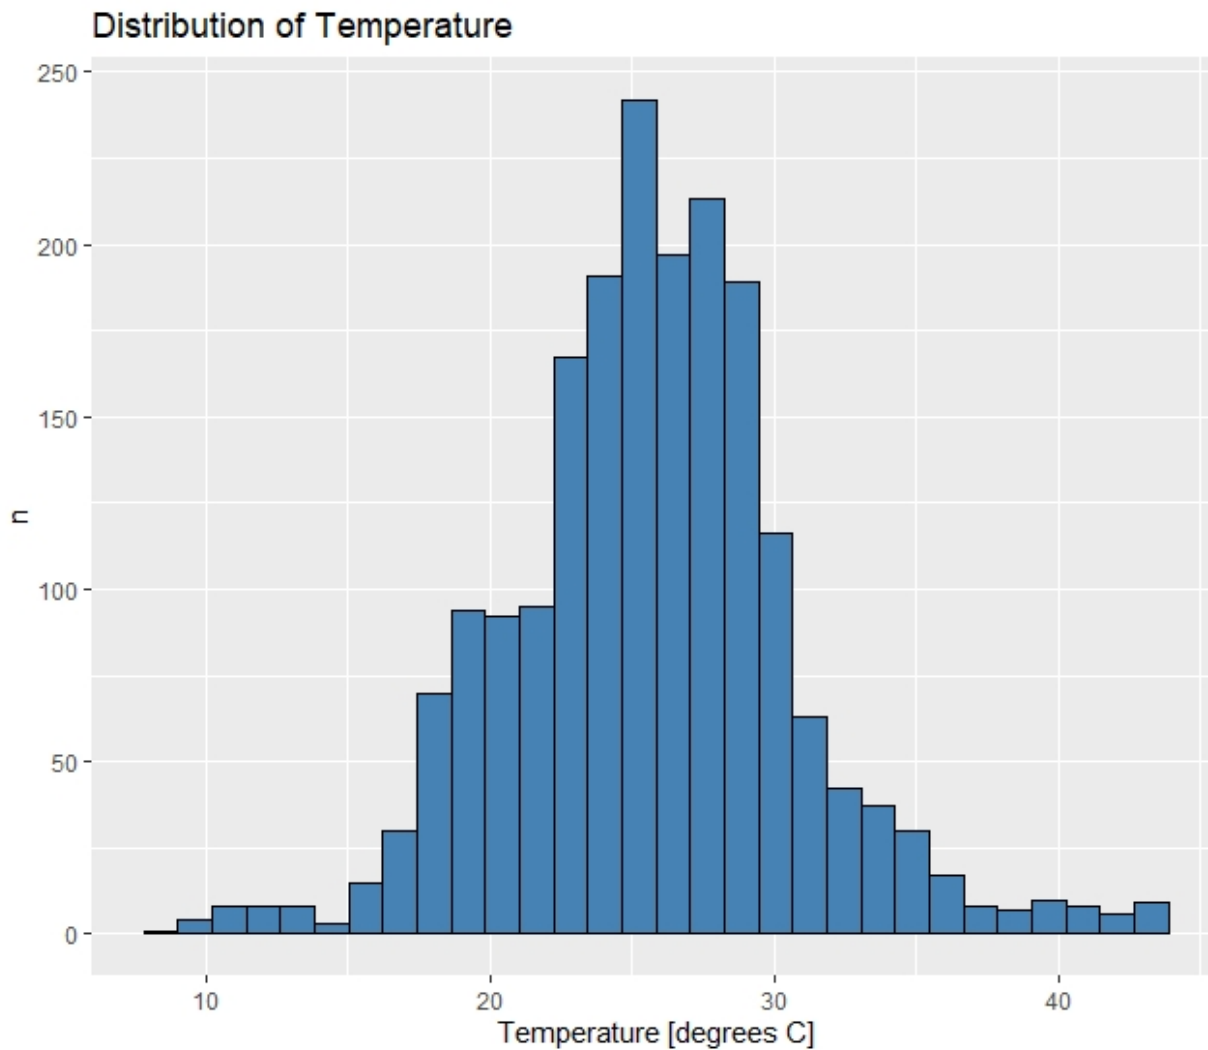
	label	mean(temperature)
1	apple	22.58299
2	banana	27.36189
3	blackgram	30.02464
4	chickpea	18.86139
5	coconut	27.50782
6	coffee	25.61020
7	cotton	23.88927
8	grapes	24.51606
9	jute	24.92283
10	kidneybeans	20.09896
11	lentil	24.45032
12	maize	22.58433
13	mango	31.16318
14	mothbeans	28.18198
15	mungbean	28.53177
16	muskmelon	28.65102
17	orange	23.17387
18	papaya	33.54618
19	pigeonpeas	28.11628
20	pomegranate	21.70613
21	rice	23.63521
22	watermelon	25.63192

Temperature for optimal crops



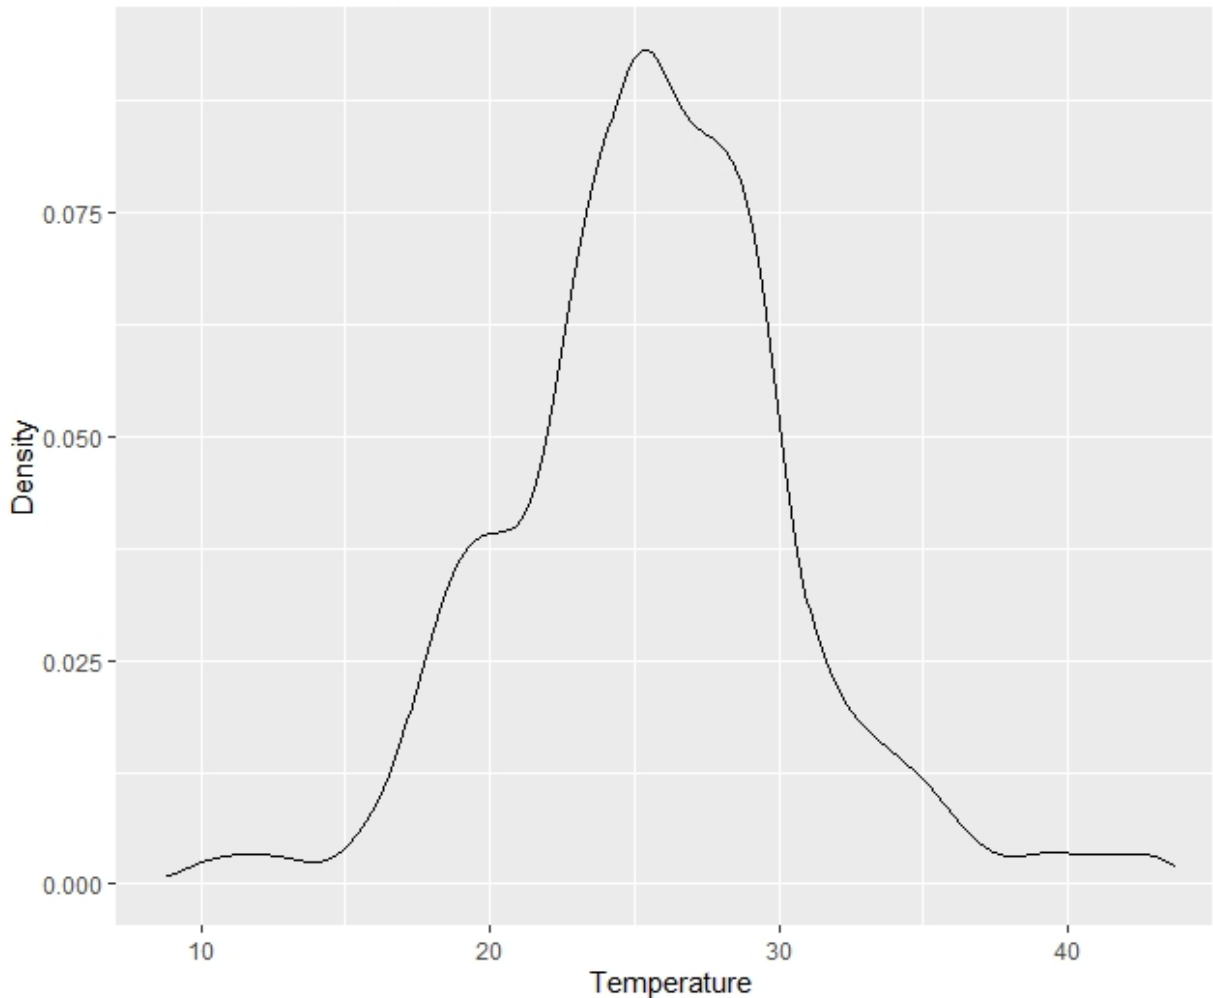

```
train %>% group_by(label) %>% ggplot(aes(x=reorder(label,temperature/m),y=temperature/m)) + geom_bar(stat = "identity", fill="steelblue") +labs(x="Crops", y="Temperature [degrees C]", title="Temperature for optimal crops") + theme(text = element_text(size=15))+coord_flip()
```

The Histogram and Density chart shows the distribution of Temperature. Also, the Temperature range is normal distributed with a mean temperature of 25.67 degrees Celsius.



```
ggplot(train,aes(temperature)) +geom_histogram(bins=30, fill="steelblue",color="black")+labs(x="Temperature [degrees C]",y="n")+ggtitle("Distribution of Temperature")
```

Distribution of Temperature



```
ggplot(train,aes(temperature)) +geom_density()+labs(x="Temperature",y="Density")+ggtitle("Distri  
bution of Temperature")
```

Rainfall

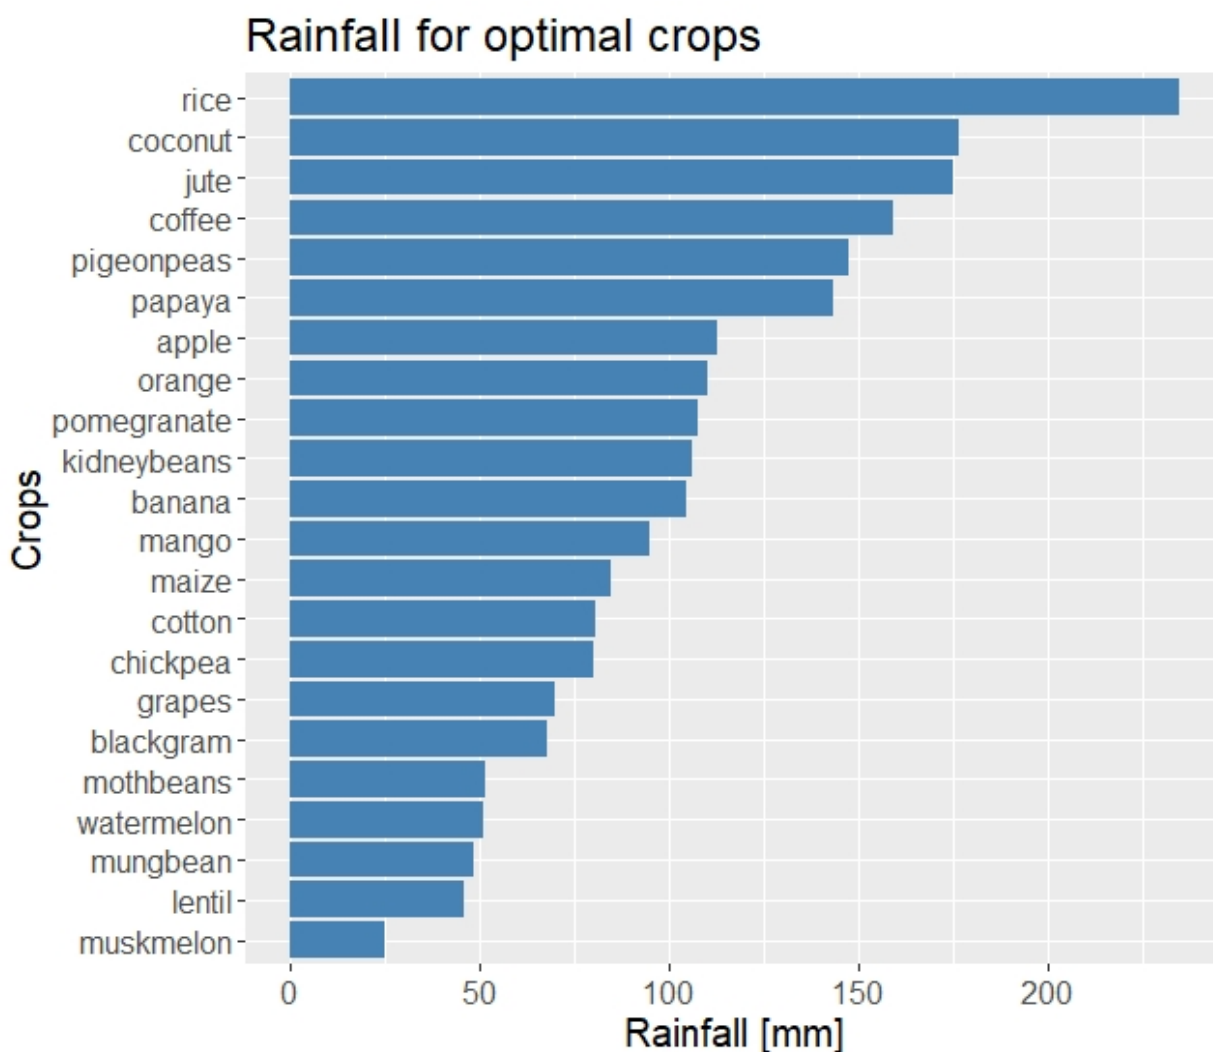
Water helps a plant by transporting important nutrients through the plant. Nutrients are drawn from the soil and used by the plant. Without enough water in the cells, the plant will droop, so water helps a plant to stand upright. Water carries dissolved sugar and other nutrients through the plant.

The water demand of the analyzed plants is very different. Rice, jute and coconut need the most water and muskmelon needs about 10-fold less water than rice. A list of all mean rain amount values is summarized in the following table and bar chart:

```
train %>%group_by(label)%>% summarise(mean(rainfall)) %>% print.data.frame()
```

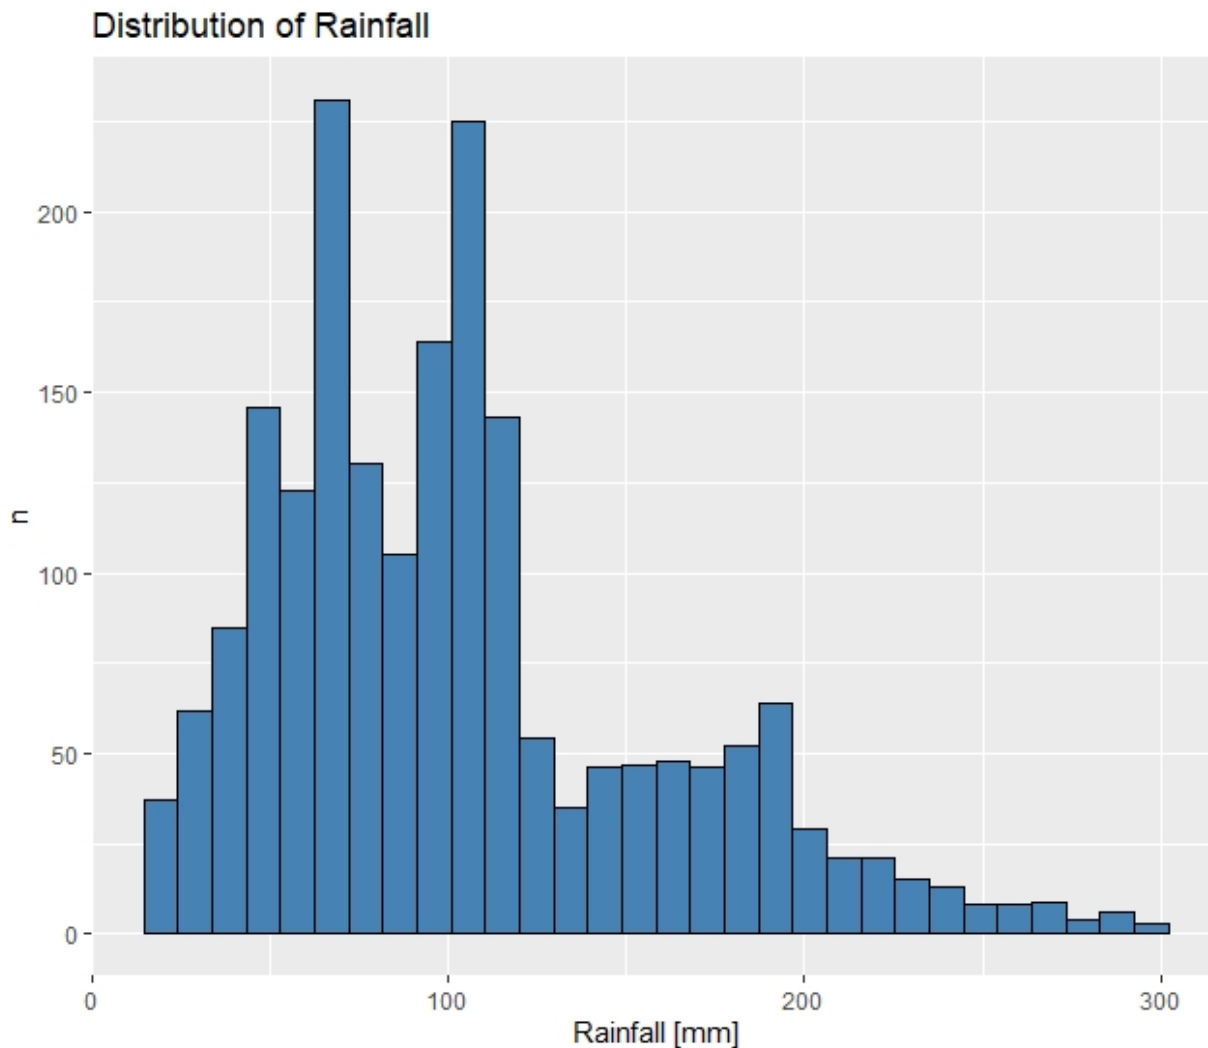
	label	mean(rainfall)
1	apple	112.94575
2	banana	104.70680
3	blackgram	67.77173
4	chickpea	79.99625
5	coconut	176.35874
6	coffee	159.14124
7	cotton	80.52756
8	grapes	69.69563
9	jute	174.78475
10	kidneybeans	105.92832
11	lentil	46.06164
12	maize	84.76890
13	mango	94.66307
14	mothbeans	51.55473
15	mungbean	48.27146
16	muskmelon	24.77892
17	orange	110.14881
18	papaya	143.62198
19	pigeonpeas	147.60782
20	pomegranate	107.62403
21	rice	234.83100
22	watermelon	50.98561

The histogram shows the different water demands.



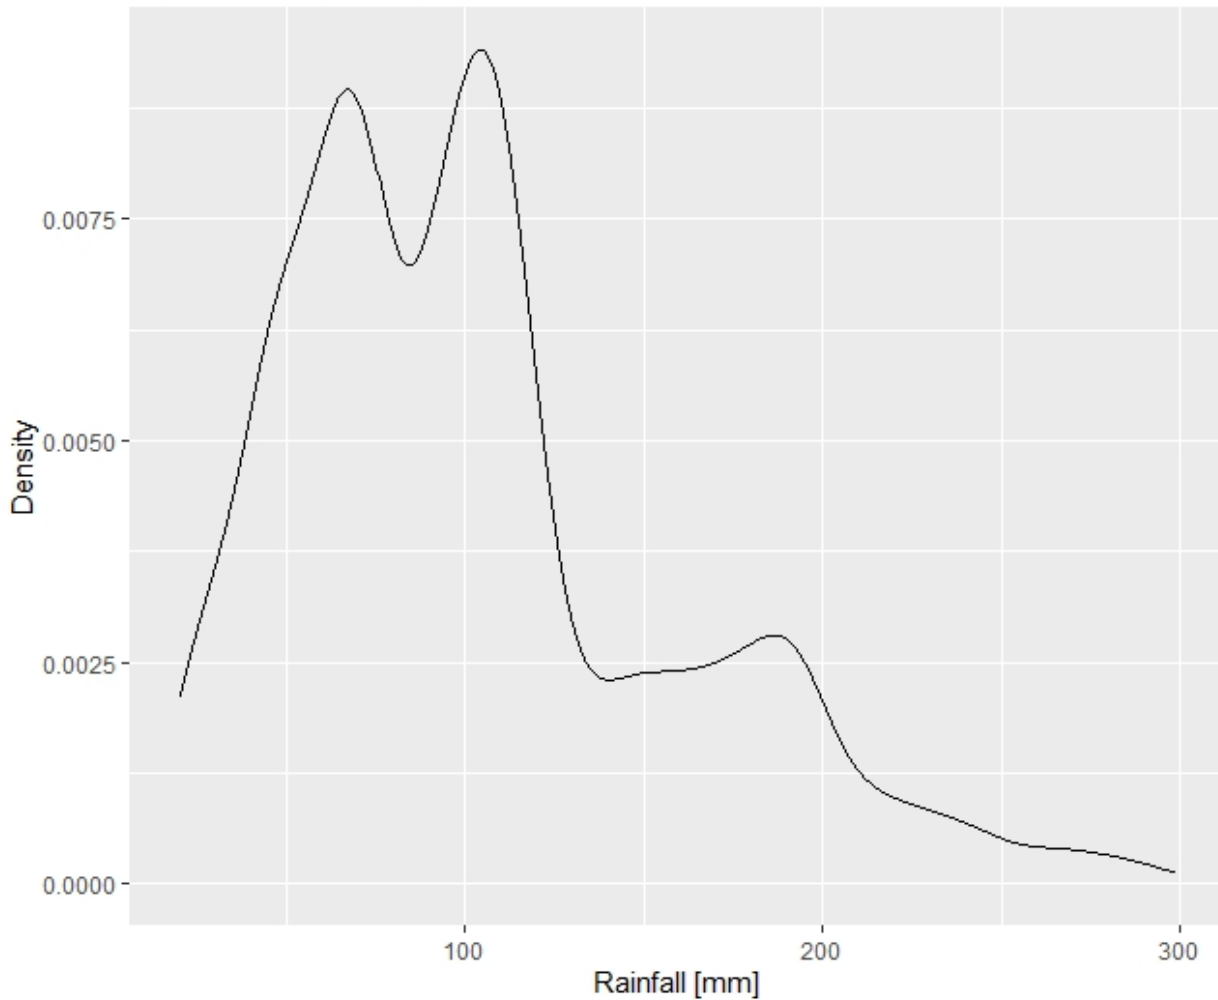
```
train %>% group_by(label) %>% ggplot(aes(x=reorder(label,rainfall/m),y=rainfall/m)) + geom_bar(stat = "identity", fill="steelblue") +labs(x="Crops", y="Rainfall [mm]", title="Rainfall for optimal crops") + theme(text = element_text(size=15))+coord_flip()
```

The Histogram and Density chart shows the distribution of rainfall amount. Individual crops require large amounts of water, but most plants need around 90mm of rainfall for optimal growth.



```
ggplot(train,aes(rainfall)) +geom_histogram(bins=30, fill="steelblue",color="black")+labs(x="Rainfall [mm]",y="n")+ggtitle("Distribution of Rainfall")
```

Distribution of Rainfall



```
ggplot(train,aes(rainfall)) +geom_density()+labs(x="Rainfall [mm]",y="Density")+ggtitle("Distribution of Rainfall")
```

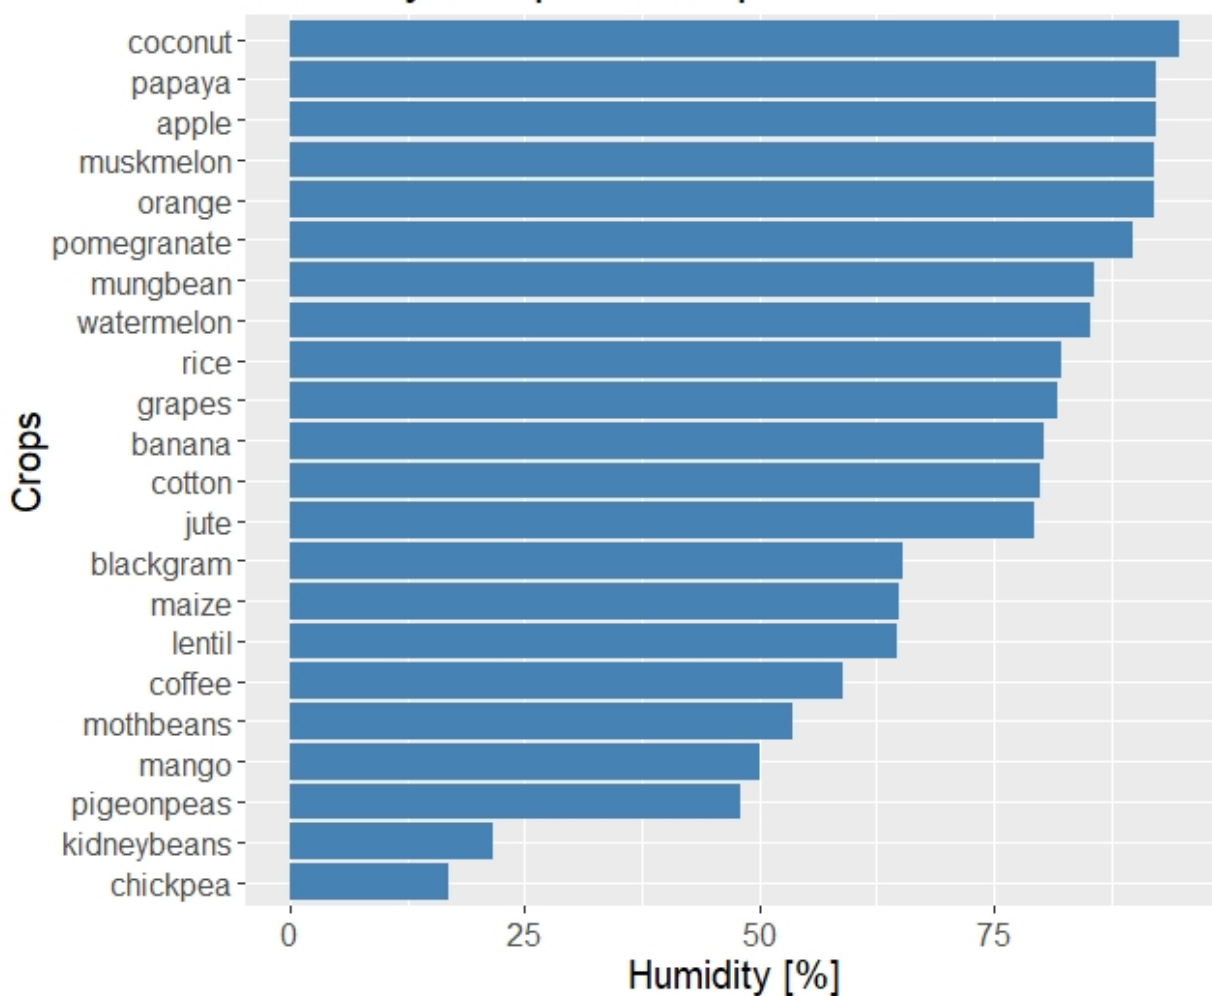
Humidity

When conditions are too humid, it may promote the growth of mold and bacteria that cause plants to die and crops to fail, as well as conditions like root or crown rot. Humid conditions also invite the presence of pests, such as fungus gnats, whose larva feed on plant roots and thrive in moist soil. Coconut, papaya and apple need the highest humidity and chickpea and kidneybeans need the lowest value. A list of all mean percent of humidity is summarized in the following table and bar chart:

```
train %>%group_by(label)%>%summarise(mean(humidity)) %>% print.data.frame()
```

	label	mean(humidity)
1	apple	92.39294
2	banana	80.42926
3	blackgram	65.34707
4	chickpea	16.89476
5	coconut	94.83994
6	coffee	58.89315
7	cotton	79.96077
8	grapes	81.85577
9	jute	79.35092
10	kidneybeans	21.56208
11	lentil	64.64839
12	maize	64.99231
13	mango	50.08699
14	mothbeans	53.67012
15	mungbean	85.65358
16	muskmelon	92.21952
17	orange	92.21013
18	papaya	92.39890
19	pigeonpeas	48.04678
20	pomegranate	89.92106
21	rice	82.30309
22	watermelon	85.31541

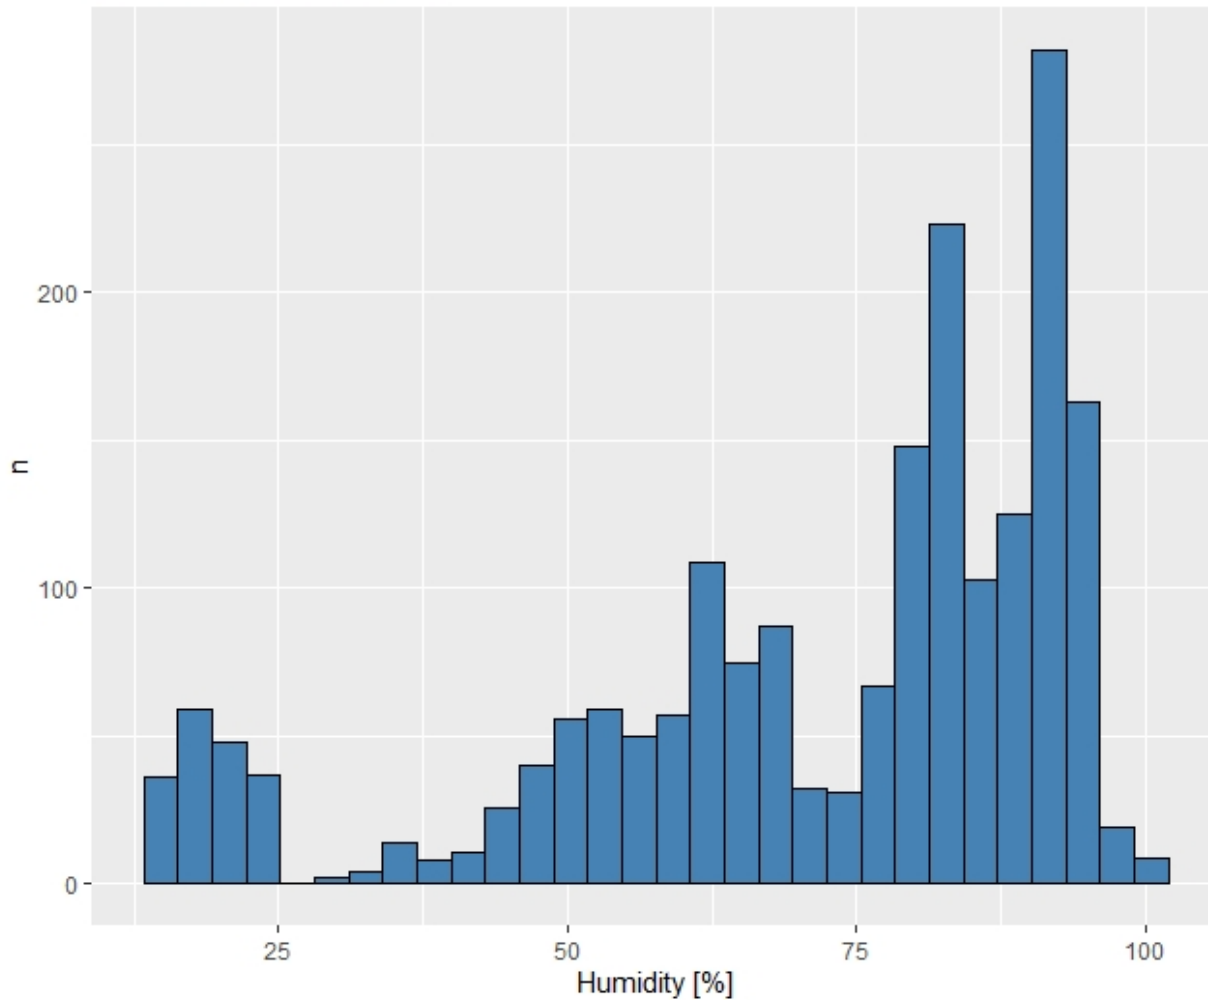
Humidity for optimal crops



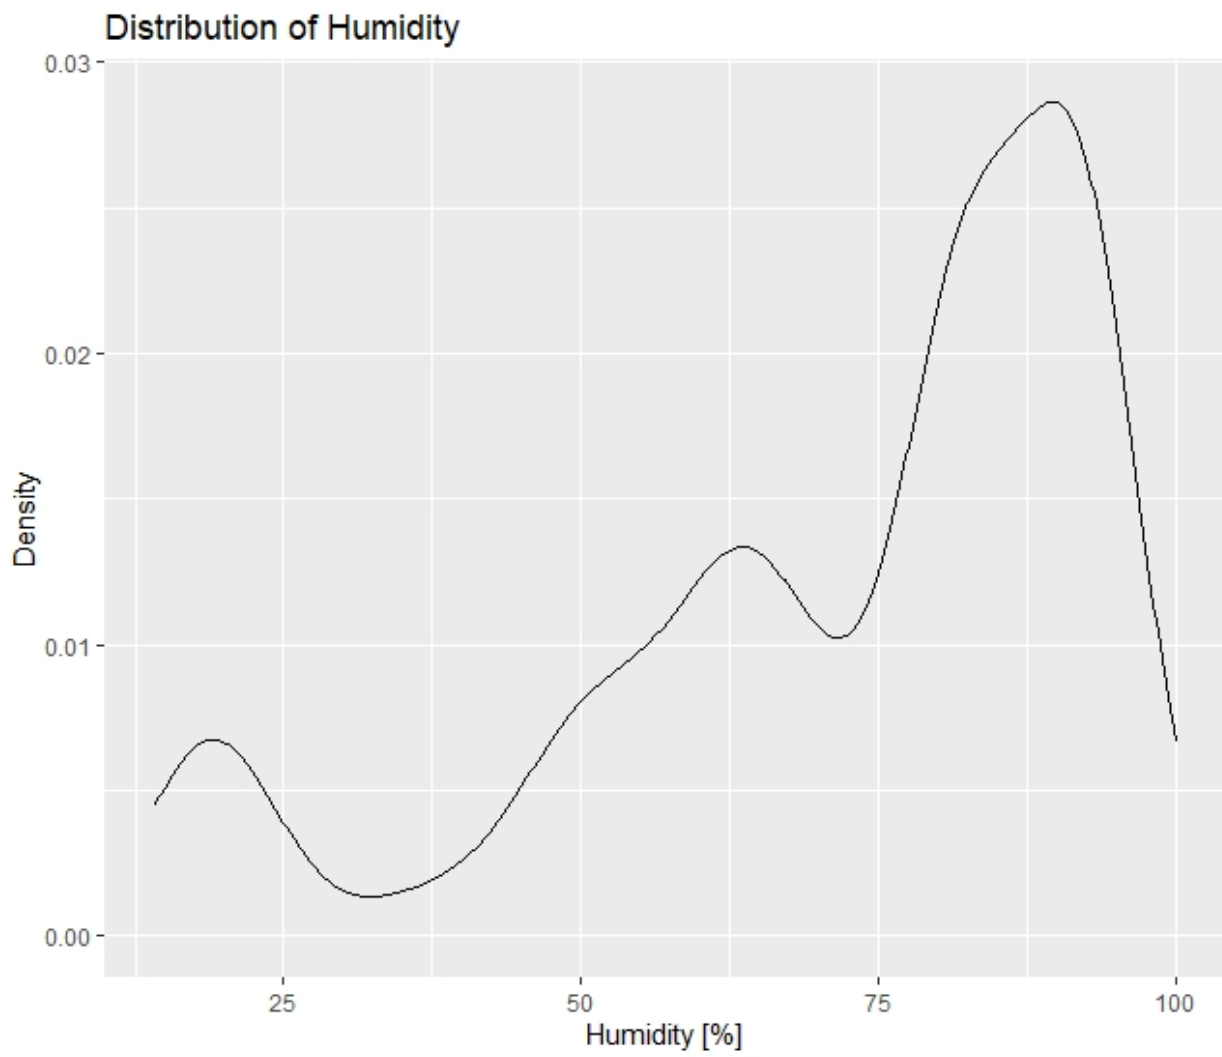
```
train %>% group_by(label) %>% ggplot(aes(x=reorder(label,humidity/m),y=humidity/m)) + geom_bar(stat = "identity", fill="steelblue") +labs(x="Crops", y="Humidity [%]", title="Humidity for optimal crops") + theme(text = element_text(size=15))+coord_flip()
```

The Histogram and Density chart shows the distribution of Humidity. The higher the humidity, the more plants can grow optimally.

Distribution of Humidity



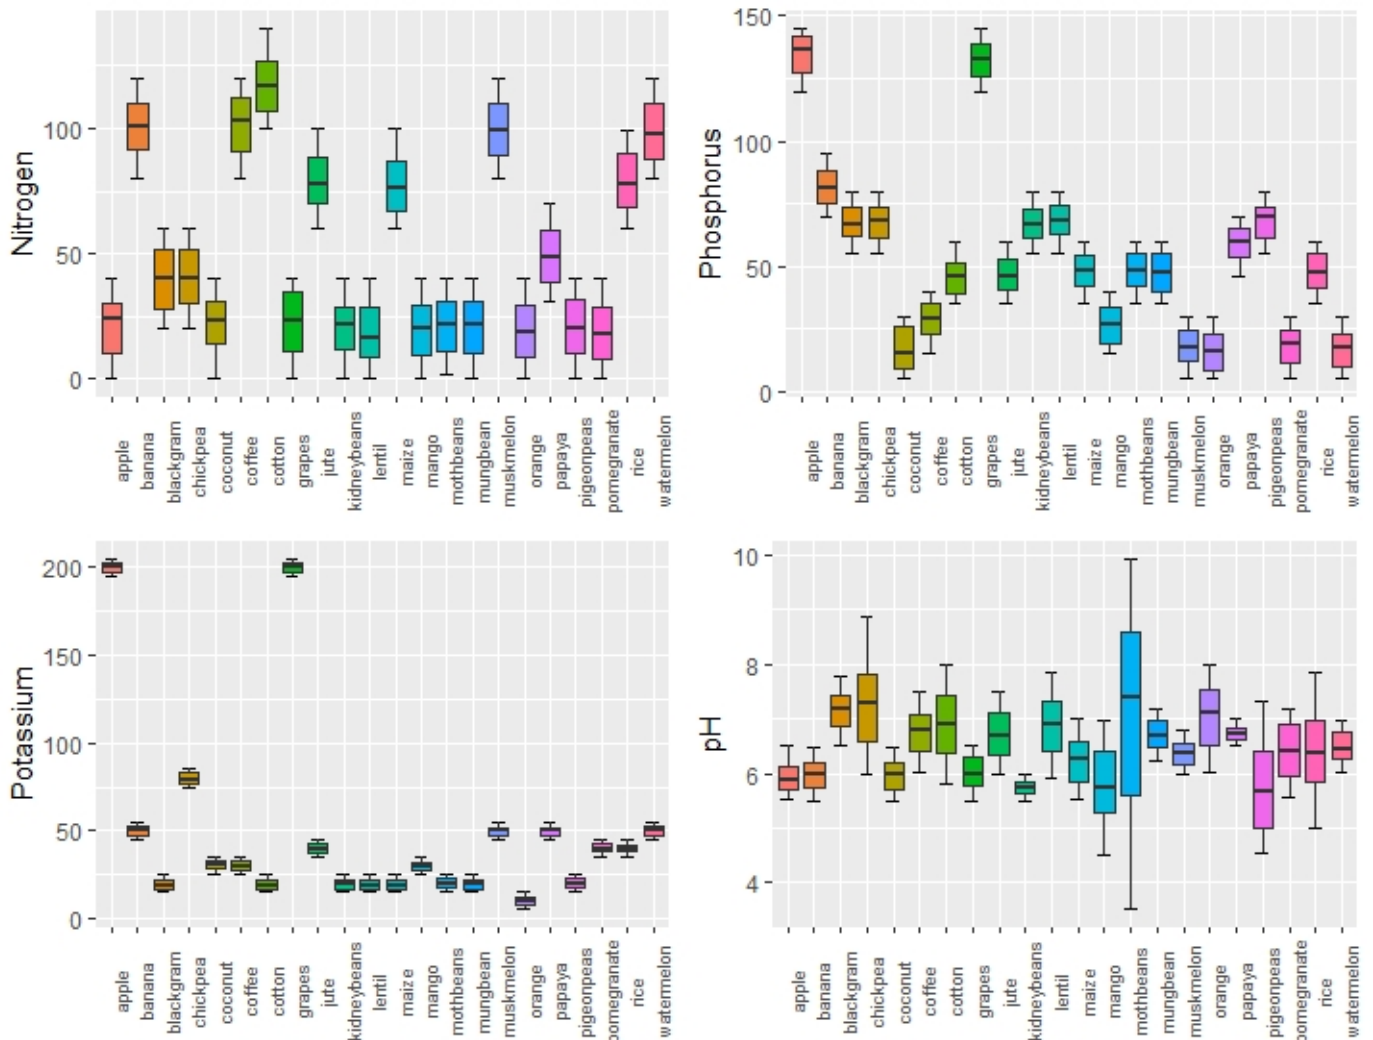
```
ggplot(train,aes(humidity)) +geom_histogram(bins=30, fill="steelblue",color="black")+labs(x="Humidity [%]",y="n")+ggtitle("Distribution of Humidity")
```



```
ggplot(train,aes(humidity)) +geom_density()+labs(x="Humidity [%]",y="Density")+ggtitle("Distribution of Humidity")
```


SOIL CONDITION RANGE

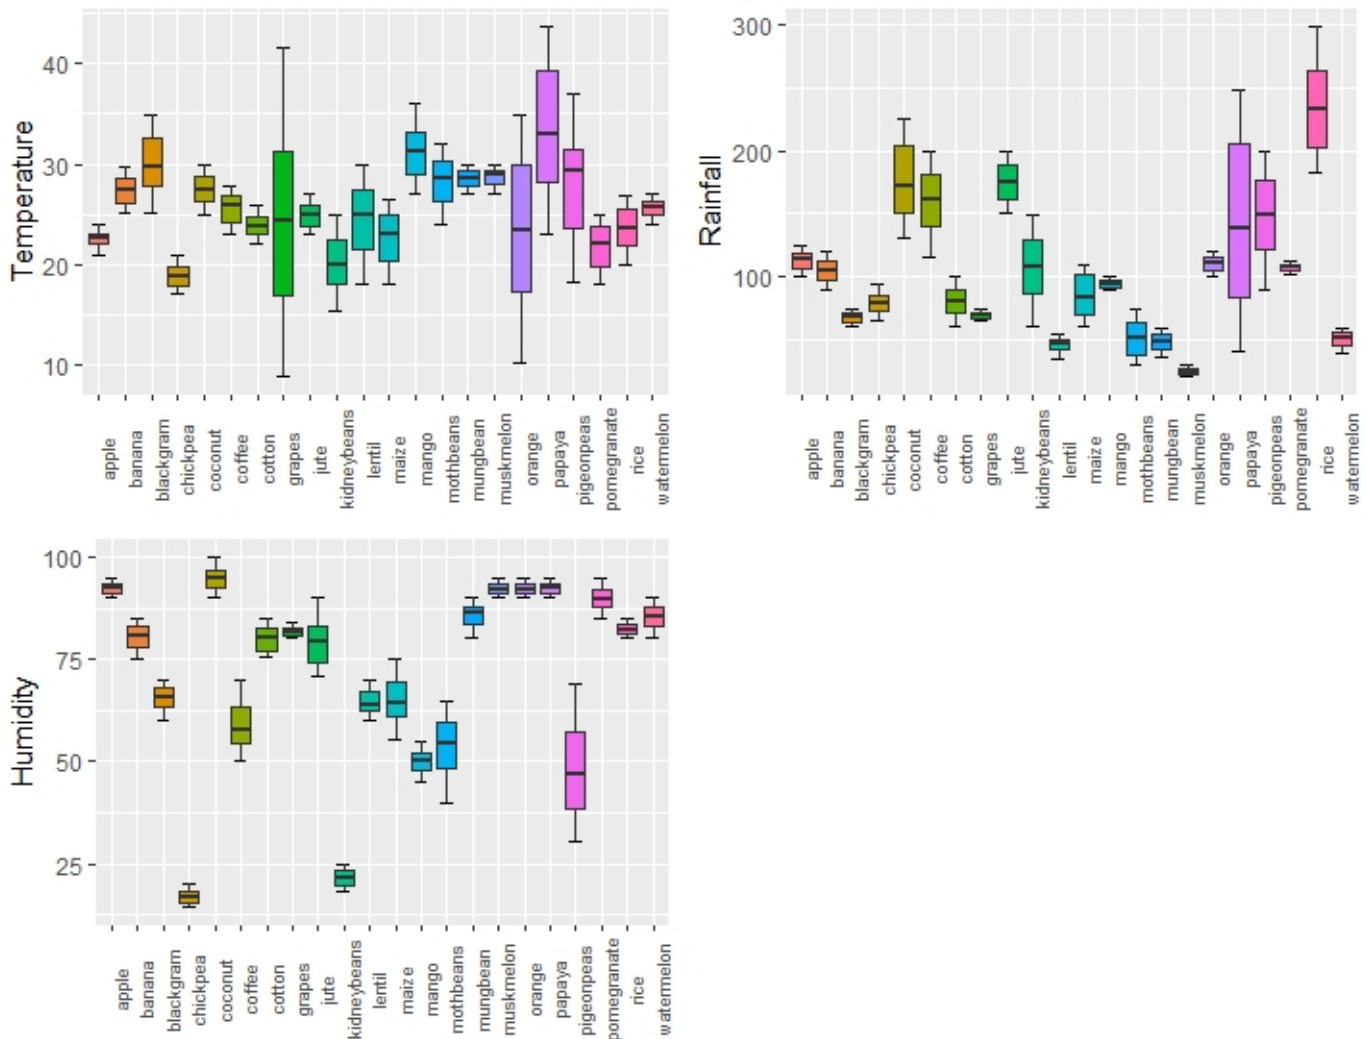
The range of the soil condition Nitrogen, Phosphorus, Potassium and pH are displayed in a following boxplots. For Nitrogen, Phosphorus and Potassium the value ranges are similar between all crops. Compared to the parameters, it is clear that there is only a small range of potassium content for optimal cultivation and harvest. The pH values vary much more among themselves. Very noticeable is the pH range of mothbeans, because it covers the whole range (3.5 - 10).



```
box1 <- ggplot(train,aes(y=N, group=label, x=label, fill=label))+stat_boxplot(geom="errorbar",width=0.5)+geom_boxplot()+ labs(y="Nitrogen") + theme (axis.text.x=element_text(angle=90, size=7), legend.position="none", axis.title.x=element_blank())
box2 <- ggplot(train,aes(y=K, group=label, x=label, fill=label))+stat_boxplot(geom="errorbar",width=0.5)+geom_boxplot()+ labs(y="Potassium") + theme (axis.text.x=element_text(angle=90, size=7),legend.position="none", axis.title.x=element_blank())
box3 <- ggplot(train,aes(y=P, group=label, x=label, fill=label))+stat_boxplot(geom="errorbar",width=0.5)+geom_boxplot()+ labs(y="Phosphorus") + theme (axis.text.x=element_text(angle=90, size=7),legend.position="none", axis.title.x=element_blank())
box4 <- ggplot(train,aes(y=ph, group=label, x=label, fill=label))+stat_boxplot(geom="errorbar",width=0.5)+geom_boxplot()+ labs(y="pH") + theme (axis.text.x=element_text(angle=90, size=7),legend.position="none", axis.title.x=element_blank())
grid.arrange(box1,box3,box2,box4, nrow=2)
```

ENVIRONMENT CONDITION RANGE

The ranges of environmental parameters vary much more than the soil parameters. This means that crops with a big parameter range are less sensitive and lead more often to optimal harvest. It also shows that environmental parameters have more influence on the optimal cultivation and harvest of crops. The biggest temperature ranges are for grapes, oranges, papaya and pigeonpeas. That means that these crops are easier to cultivate. Rainfall sensitivity is lowest for papaya, pigeonpeas, rice and coconut. Humidity sensitivity is lowest for pigeonpeas. In summary, pigeonpeas can be grown in a wide range of environmental parameters and can lead to an optimal harvest.



```
box5 <- ggplot(train,aes(y=temperature, group=label, x=label, fill=label))+stat_boxplot(geom="errorbar",width=0.5)+geom_boxplot()+ labs(y="Temperature") + theme (axis.text.x=element_text(angle=90, size=7),legend.position="none", axis.title.x=element_blank())
box6 <- ggplot(train,aes(y=rainfall, group=label, x=label, fill=label))+stat_boxplot(geom="errorbar",width=0.5)+geom_boxplot()+ labs(y="Rainfall") + theme (axis.text.x=element_text(angle=90, size=7),legend.position="none", axis.title.x=element_blank())
box7 <- ggplot(train,aes(y=humidity, group=label, x=label, fill=label))+stat_boxplot(geom="errorbar",width=0.5)+geom_boxplot()+ labs(y="Humidity") + theme (axis.text.x=element_text(angle=90, size=7),legend.position="none", axis.title.x=element_blank())
grid.arrange(box5,box6,box7, nrow=2)
```

4. RESULTS - MACHINE LEARNING ALGORITHM

To see how this is a type of machine learning, we need to build an algorithm with data we have collected. as users look for crop recommendations. The test set for this calculation of accuracy is provided in the data that will then be applied outside our control, as farmer look for crop recommendations. Lets create a test set to assess the accuracy of the models we implement.

LOSS FUNCTION

We define the rating for movie i by user and denote our prediction. The residual mean squared error RMSE is defined as:

$$RMSE < -function(true_{ratings}, predicted_{ratings})sqrt(mean((true_{ratings} - predicted_{ratings})^2))$$

The lower the RMSE, the better it is.

4.1. FIRST MODEL

We start to build the simplest possible recommendation system. We predict the Nitrogen on the crops regardless of the label. The estimation minimizes the RMSE is the least squares estimation of the rating and is the average of all ratings.

```
mu <- mean(train$temperature)
mu

[1] 25.67037
```

For the prediction of all unknown ratings with $\hat{\mu}$ following RMSE will be calculated:

```
naive_rmse <- RMSE(validation$temperature, mu)
naive_rmse

[1] 5.427997
```

The results will be saved in a table.

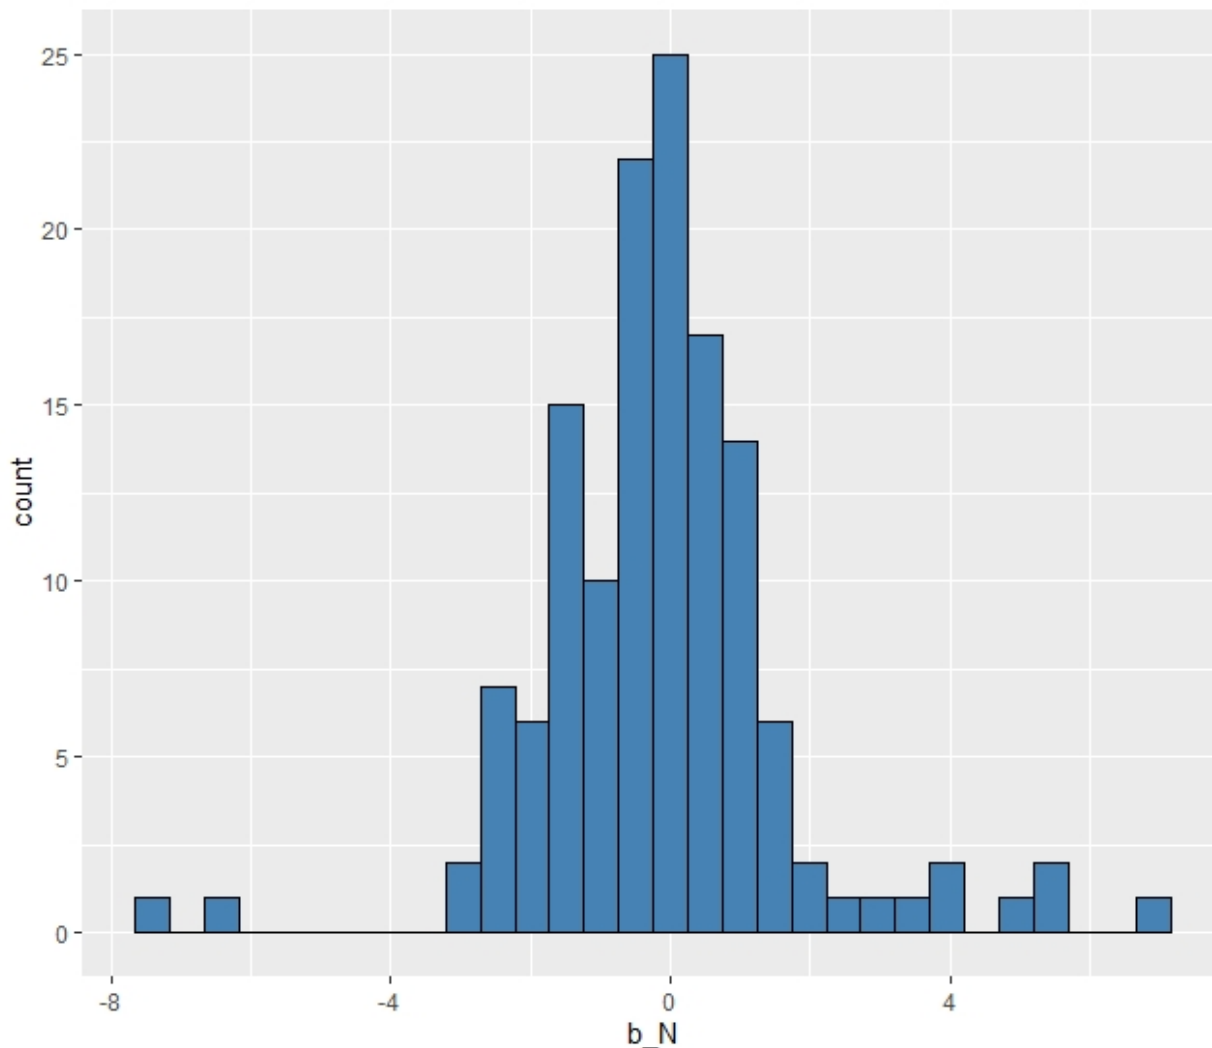
```
rmse_results <- data_frame (method="Just the average", RMSE=naive_rmse)
rmse_results
```

4.2. Nitrogen Effect

The `lm()` function will be very slow and so I compute it in the whole model using the average like in the following way:

```
mu <- mean(train$temperature)
N_avgs <- train %>% group_by(N) %>% summarise(b_N=mean(temperature-mu))
```

We can see that these estimates vary substantially:



```
train %>% group_by(N) %>% summarise(b_N=mean(temperature-mu)) %>% ggplot(aes(b_N)) +geom_histogram(bins=10, color="black", fill="steelblue")
```

We can see how much our prediction improves once we use following calculation: NAs in the data will be changed to Zero, otherwise RMSE could not be calculated.

```
predicted_temperature <- mu + validation %>%left_join(N_avgs, by="N") %>% pull(b_N)
predicted_temperature[is.na(predicted_temperature)]<- 0
N_effect_rmse <- RMSE(validation$temperature,predicted_temperature)
N_effect_rmse

[1] 5.828109
```

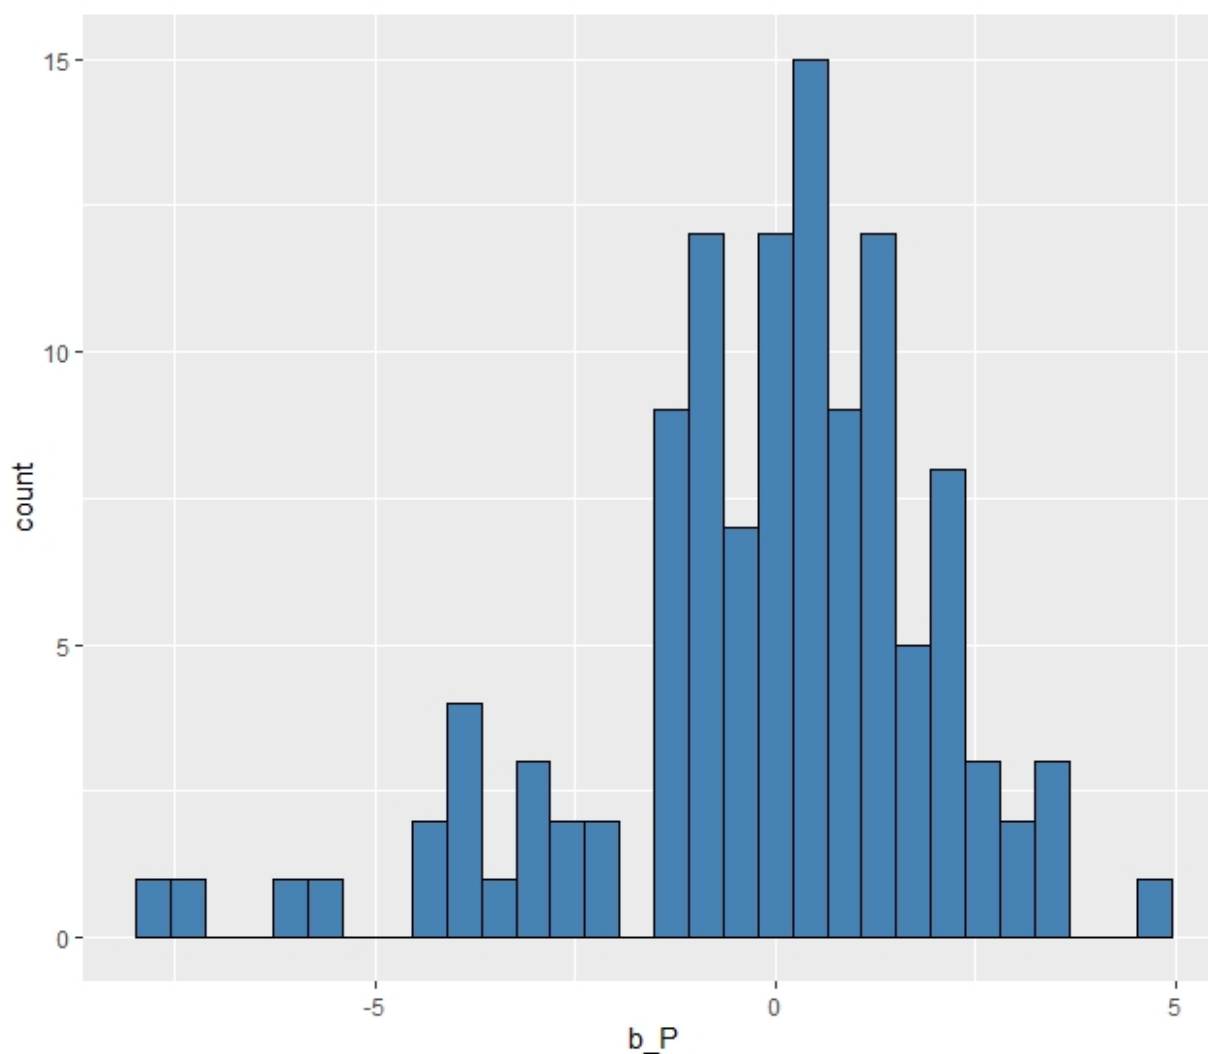
By the Nitrogen Effect the RMSE increase. The results will be saved in the table:

```
rmse_results <- bind_rows(rmse_results, data_frame(method="Nitrogen Effect Model", RMSE=N_effect_rmse))
rmse_results
```

4.3. PHOSPOURUS EFFECT

The approximation is computed by the phosphorus average, and we can construct predictors and see how much the RMSE improves:

```
mu <- mean(train$temperature)
P_avgs <- train %>% group_by(P) %>% left_join(N_avgs, by="N") %>% summarise(b_P=mean(temperature-mu-b_N))
```



```
train %>% group_by(P) %>% summarise(b_P=mean(temperature-mu)) %>% ggplot(aes(b_P)) +geom_histogram(bins=10, color="black", fill="steelblue")
```

We can compute the predictors and see how the RMSE improve again. NAs in the data will be changed to Zero, otherwise RMSE could not be calculated.

```

predicted_temperature <- validation %>%left_join(N_avgs, by="N") %>% left_join(P_avgs, by="P") %
>% mutate (pred=mu+b_N +b_P)%>% pull(pred)
predicted_temperature[is.na(predicted_temperature)]<- 0

P_effect_rmse <- RMSE(predicted_temperature,validation$temperature)
P_effect_rmse

[1] 6.006829

```

The RMSE increase further. The results will be saved in a table:

```

rmse_results <- bind_rows(rmse_results, data_frame(method="Phosphorus Effect Model", RMSE=P_effect_rmse))
rmse_results

```

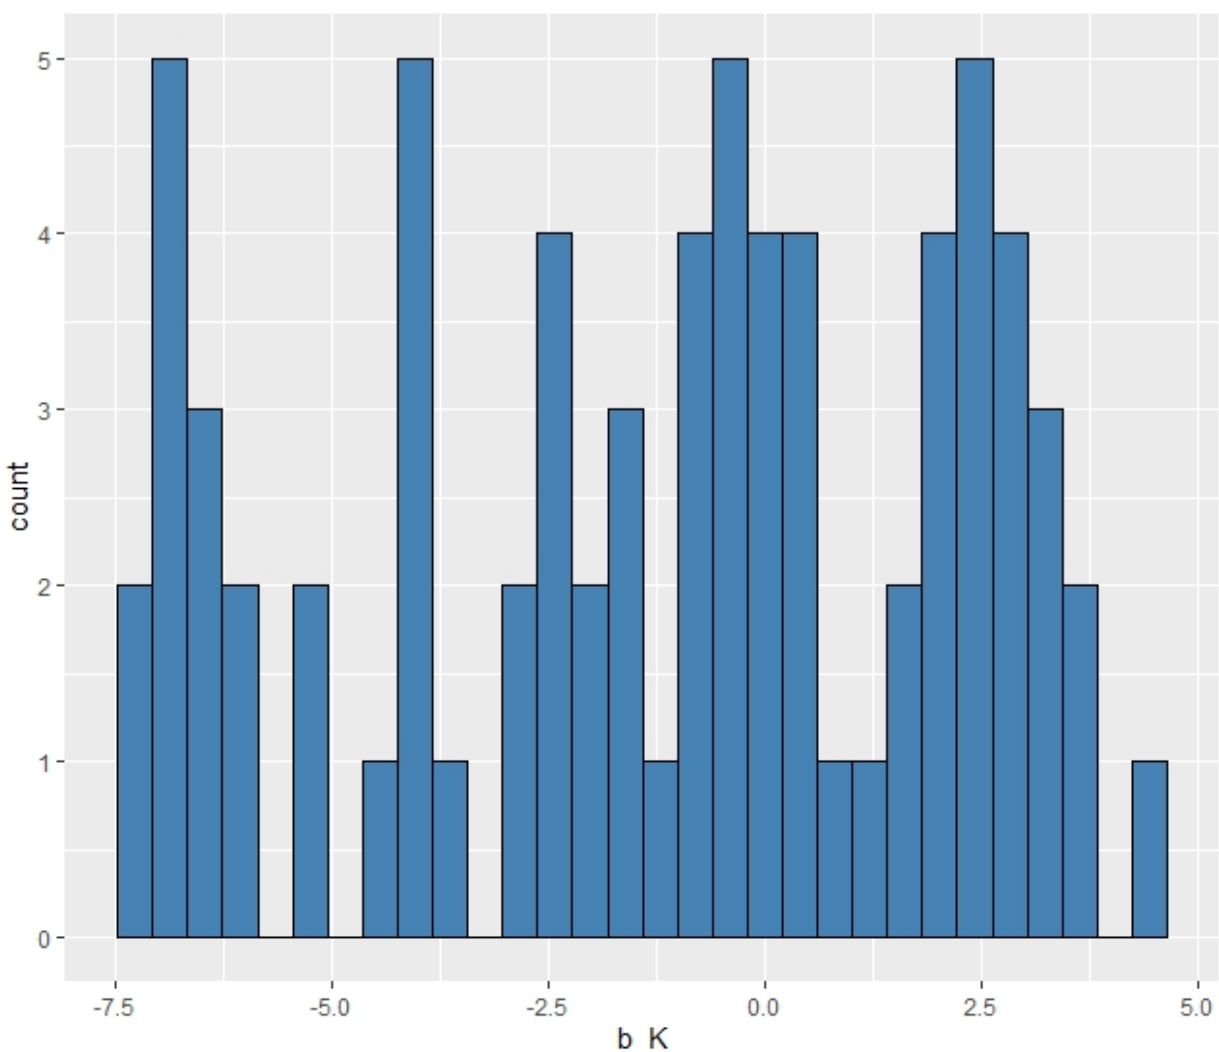
4.4. POTASSIUM EFFECT

The approximation is computed by the Potassium average, and we can construct predictors and see how much the RMSE improves:

```

mu<- mean(train$temperature)
K_avgs <- train %>% group_by(K) %>% left_join(N_avgs, by="N") %>% left_join(P_avgs, by="P")%>%
summarise(b_K=mean(temperature-mu-b_N-b_P))

```



```
train %>% group_by (K) %>% summarise(b_K=mean(temperature-mu))%>% ggplot(aes(b_K)) +geom_histogram(bins=30, color="black", fill="steelblue")
```

We can compute the predictors and see how the RMSE improve again:

```
predicted_temperature <- validation %>%left_join(K_avgs, by="K") %>% left_join(N_avgs, by="N")%>%left_join(P_avgs, by="P") %>% mutate (pred=mu+b_N +b_P+b_K)%>% pull(pred)
predicted_temperature[is.na(predicted_temperature)]<- 0

K_effect_rmse <- RMSE(predicted_temperature,validation$temperature)
K_effect_rmse

[1] 5.623298
```

The Potassium Effect let the RMSE decrease. The results will be saved in a table:

```
rmse_results <- bind_rows(rmse_results, data_frame(method="Potassium Effect Model", RMSE=K_effect_rmse))
rmse_results
```

4.5. CHOOSING THE PENALTY TERMS

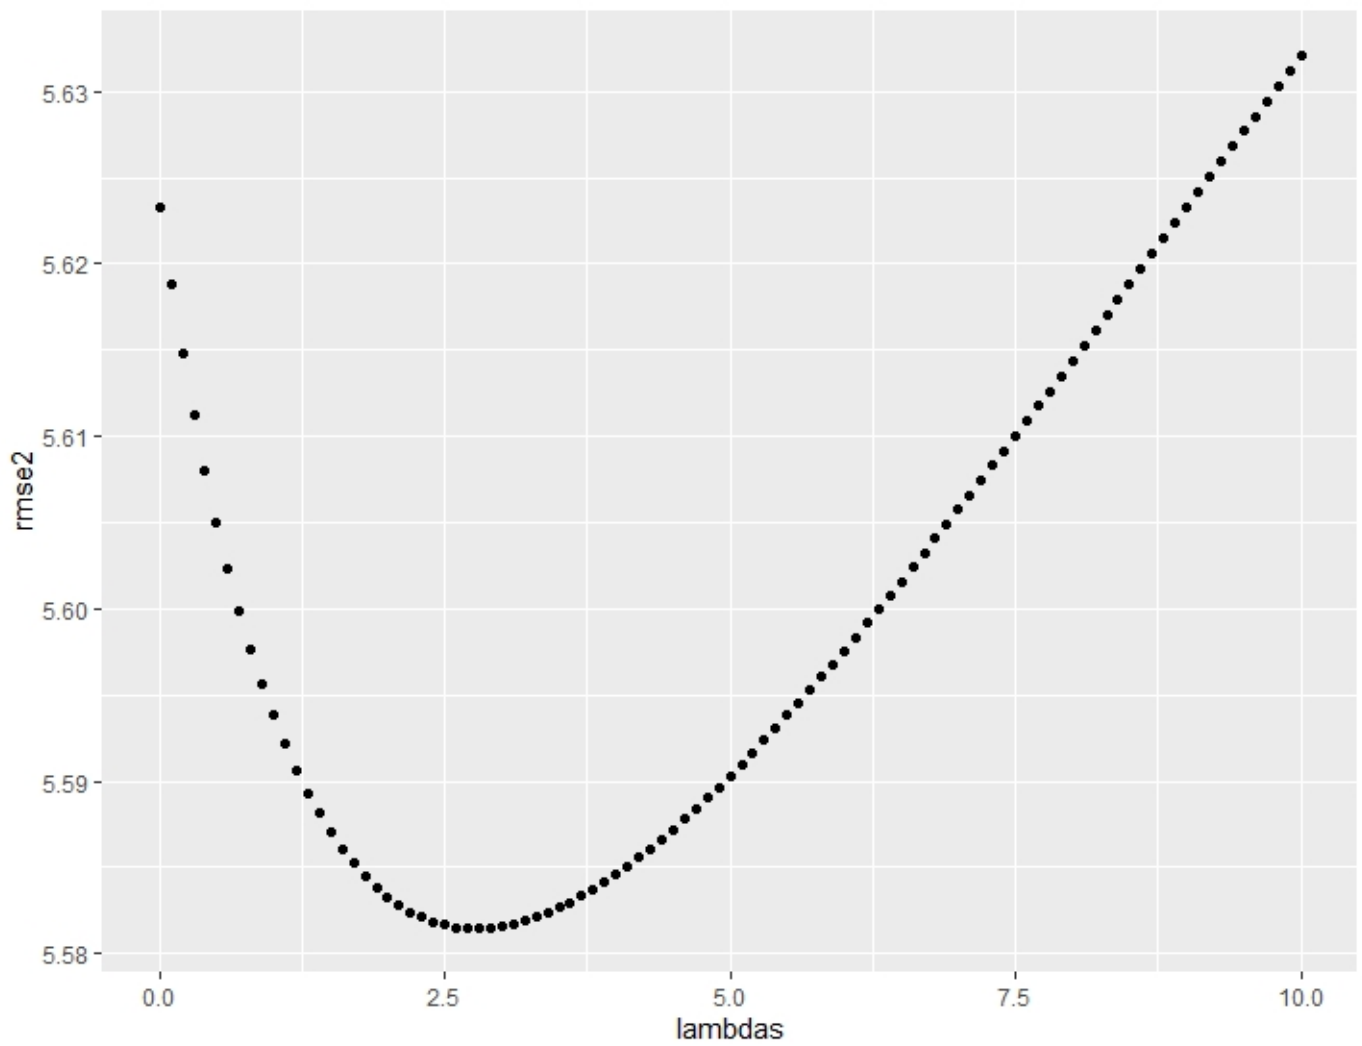
Despite the large variation of the different factors our improvements in RMSE is good. If n is very large, the estimation will be very stable, and the penalty λ is effectively ignored since $n_i + \lambda$ is about n_i . When the n_i is small, then the estimation $\hat{b}_i(\lambda)$ is shrunk towards 0. The larger λ , the more we shrink. We can optimize this model by using λ . Lets compute these regularized estimate of b_N , b_P and b_K with λ which is a turning parameter and a sequence between 0 and 10 with distances of 0.25 will be generate and compute the lowest RMSE.

```
lambdas <-seq(0,10,0.1)
mu<- mean(train$temperature)

rmse2 <- sapply(lambdas,function(l){
  mu<-mean(train$temperature)
  b_N_reg<-train %>% group_by(N)%>% summarise(b_N_reg=sum(temperature-mu)/(n()+1))
  b_P_reg<-train %>% group_by(P) %>% left_join(b_N_reg, by="N") %>% summarise(b_P_reg=sum(temperature-mu-b_N_reg)/(n()+1))
  b_K_reg<-train %>% group_by(K) %>% left_join(b_N_reg, by="N") %>% left_join(b_P_reg, by="P") %>% summarise(b_K_reg=sum(temperature-mu-b_N_reg-b_P_reg)/(n()+1))

  predicted_temperature <- validation %>% left_join(b_N_reg, by="N")%>% left_join(b_P_reg, by="P")%>% left_join(b_K_reg, by="K")%>% mutate(pred=mu+b_N_reg+b_P_reg+b_K_reg) %>% pull(pred)
  predicted_temperature[is.na(predicted_temperature)]<- 0
  return(RMSE(predicted_temperature, validation$temperature)))

qplot(lambdas,rmse2)
```



```
lambda_low<- lambdas[which.min(rmse2)]  
lambda_low  
  
[1] 2.8
```

The optimal lambda with the lowest RMSE is not calculable. The new RMSE for the model with the optimal lambda is compute as:


```

mu<- mean(train$temperature)
N_reg<-train %>% group_by(N)%>% summarise(N_reg=sum(temperature-mu)/(n()+lambda_low))
P_reg<-train %>% group_by(P) %>% left_join(N_reg, by="N") %>% summarise(P_reg=sum(temperature-mu
-N_reg)/(n()+lambda_low))
K_reg<-train %>% group_by(K) %>% left_join(N_reg, by="N") %>% left_join(P_reg, by="P") %>% summa
rise(K_reg=sum(temperature-mu-N_reg-P_reg)/(n()+lambda_low))

predicted_temperature <- validation %>% left_join(N_reg, by="N")%>% left_join(P_reg, by="P")%
>% left_join(K_reg, by="K")%>% mutate(pred=mu+N_reg+P_reg+K_reg) %>% pull(pred)
predicted_temperature[is.na(predicted_temperature)]<- 0

All_reg_effect_rmse <- RMSE(predicted_temperature,validation$temperature)
All_reg_effect_rmse

[1] 5.581457

rmse_results <- bind_rows(rmse_results, data_frame(method="Regularized Model with optimal lambd
a", RMSE=All_reg_effect_rmse))
rmse_results

```

The summarized RMSE are displayed in the following table:

A tibble: 5 x 2

method	RMSE
<chr>	<dbl>
1 Just the average	5.43
2 Nitrogen Effect Model	5.83
3 Phosphorus Effect Model	6.01
4 Potassium Effect Model	5.62
5 Regularized Model with optimal lambda	5.58

5. CONCLUSION

The objective of this project was to analyze the crop recommendation data and develop a machine learning algorithm from the dataset "Crop recommendation" from Kaggle website. For this dataset different analyzes were done and a machine learning algorithm was calculated. To see how close the predicted value to the actual value is, the RMSE was calculated. I developed a naive model and different models with effects of rainfall, humidity, Nitrogen, Potassium, Phosphorous and pH with their regarding RMSE. This model is tested with the validation dataset. This algorithm includes an error loss, which is included in the last calculation. The optimal lambda is calculate with 2.8. The end RMSE is calculate with 5.5814 and that means that the RMSE is increased by 0.15.